

Conducting Qualitative Interviews with AI

Felix Chopra

Ingar Haaland

First version: September 15, 2023

This version: March 9, 2026

Abstract

Qualitative interviews offer rich insights into human behavior, but their lack of scalability has limited their use in economics. This paper introduces and evaluates AI-led interviews as a novel method for collecting qualitative data. Across different economic contexts, a comprehensive set of quality metrics demonstrates the high quality of AI-led interviews. We then demonstrate that AI-led interviews generate thematically richer responses than other survey methods for collecting qualitative data at scale. A decomposition exercise shows that the benefits of AI-led interviews mainly arise from asking dynamic probing questions. Importantly, AI-led interviews make themes visible that would otherwise be unobserved. These themes matter: We show that responses in AI-led interviews are internally valid and predict future behavior six months later. Finally, we demonstrate that AI-led interviews allow hypothesis generation about important economic phenomena, reaching thematic saturation in small samples comparable to studies using human-led interviews.

Keywords: Artificial Intelligence, Interviews, Large Language Models, Qualitative Methods, Mental Models, Surveys.

JEL codes: C83, C90, D14, D91, Z13

We thank Peter Andre, Elliot Ash, Luca Braghieri, Kai Barron, Leonardo Bursztyn, Benjamin Enke, Tore Ellingsen, Luca Henkel, Brian Jabarian, Rafael Jiménez-Durán, Christopher Roth, Carlo Schwarz, Sonja Settele, and Johannes Wohlfart for helpful discussions. We also thank seminar participants at Aarhus University, the Early Career Behavioral Economics Conference in Bonn, Berlin (BBE Seminar), the briq Belief Workshop, Copenhagen Business School, ETH Zurich Narratives in Economics Workshop, IZA, Johannes Kepler University Linz, Lund University, NHH Norwegian School of Economics, OsloMet, Stockholm (IFN), Toulouse School of Economics, University of Bergen, University of Cologne, University of Copenhagen, University of Innsbruck, University of Oslo, University of Vienna, and Uppsala University for very helpful comments. We thank Abbey Horner, Miguel Camacho Horvitz, Ida Ingstad, Andreas Keldke, Felix Sandblom, and Rav Seime for excellent research assistance. IRB approval was obtained from the ethics committee of NHH Norwegian School of Economics (NHH-IRB 54/23; NHH-IRB-2025-98). The activities of the Center for Economic Behavior and Inequality (CEBI) are financed by the Danish National Research Foundation, Grant DNR134. Financial support from the Research Council of Norway through its Centre of Excellence Scheme (FAIR project No. 262675) is gratefully acknowledged. Chopra: Frankfurt School of Finance & Management, CESifo, IZA (f.chopra@fs.de); Haaland: NHH Norwegian School of Economics, CEPR, CESifo (Ingar.Haaland@nhh.no).

1 Introduction

A full understanding of human behavior requires qualitative data on how people perceive and reason about the world. Economists use qualitative data from open-ended survey questions to study many decision-relevant factors that are often unobserved in traditional datasets, such as narratives and mental models (Ash et al., 2026; Andre et al., 2026, 2024; Hübsch et al., 2026), economic reasoning (Stantcheva, 2021), or attention allocation (Link et al., 2024). More broadly, qualitative data has been used to study political bias, polarization and stereotypes (Ash et al., 2022; Braghieri et al., 2024; Draca and Schwarz, 2024; Durán et al., 2025).

While qualitative data has transformed our understanding of many important topics in economics (Gentzkow et al., 2019; Haaland et al., 2025), economists rarely conduct qualitative interviews. Qualitative interviews are a leading method for collecting qualitative data in other social sciences, allowing researchers to interactively ask open-ended questions and follow up on key themes rather than being restricted to a predetermined survey structure (Knott et al., 2022). The limited adoption likely reflects many factors, such as the incompatibility of traditional interviews with large-scale surveys economists typically conduct, and the focus on quantification and hypothesis testing rather than hypothesis generation (Ludwig and Mullainathan, 2024).

Artificial intelligence (AI) promises an opportunity to integrate qualitative interviews into large-scale surveys by delegating the task of interviewing human subjects to an AI interviewer. This raises several open questions about the value of AI-led qualitative interviews as a tool for studying economic questions. First, is an AI interviewer able to conduct high-quality interviews? Second, can AI-led interviews uncover novel findings of interest to economists? Third, to what extent are responses in AI-led interviews internally and externally valid? Fourth, how do qualitative interviews compare to other survey methods for collecting qualitative data in terms of data quality? In this paper, we seek to answer these questions.

We develop the first scalable AI interviewer and collect qualitative interview data at scale to comprehensively examine the quality of AI-led interviews. Since there is no consensus among qualitative researchers on how to evaluate the quality of qualitative data, we introduce several complementary quality metrics, including participants’ self-evaluation of the interviews and a thorough human evaluation of the interviews, including an assessment of adherence to best practices for qualitative interviewing, extent of hallucinations, and alignment with human judgment. Across these metrics, we consistently find that the AI interviewer conducts high-quality interviews. Turning to hypothesis generation, we show that AI-led interviews can be used for qualitative research by demonstrating their ability to identify novel insights across two different domains of interest to economists. Importantly, AI-led interviews achieve thematic saturation already in small samples commonly used in studies with human-led interviews. We next leverage our unique large-scale interview data to address fundamental methodological questions about AI-led qualitative interviews. Specifically, we show that AI-led interview responses are internally

valid for studying economic behavior and predictive of future behavior. Finally, we directly compare AI-led interviews with other survey methods for collecting qualitative data at scale, providing evidence that AI-led interviews yield quantitatively and qualitatively richer responses compared to asking open-ended survey questions. Before presenting these results in more detail, we introduce our AI interviewing method.

Our scalable AI interviewer application prompts a frontier large language model (LLM) to dynamically generate questions for an interview with a human subject in real time. The key input for the LLM is a short prompt with detailed guidelines on interviewing derived from the literature. For example, the prompts explain that it is central to ask neutral, open-ended questions instead of narrowly framed questions that prime respondents to answer in a specific way. The core prompts do not include application-specific instructions, making them portable across contexts. To conduct interviews on a new topic, researchers have to write down a short *topic guide* that specifies the length and scope of the topics that the AI interviewer should ask about, giving the researcher full control over the interview structure. Our application is modular and includes many optional features, such as (i) content moderation of interviews to detect problematic messages (Jiménez-Durán, 2021), (ii) generating on-the-fly analytic memos to help the LLM focus its attention on important details in very long interviews, and (iii) having separate LLMs for different types of interview questions to benefit from task separation (Khot et al., 2023). Participants interact with the AI interviewer through a chat interface that can be seamlessly embedded into standard survey software.¹ The chat interface is programmatically protected against participants inserting AI-generated text. To answer an interview question, participants can either type their response or record a voice message that we automatically transcribe. Within a second, the next interview question is then shown to participants. Our application is open-source with production-ready code that can be deployed at zero cost within the free tier of standard cloud infrastructure.² With LLM usage costs below \$0.10 per interview hour, our method enables researchers to run unlimited interviews in parallel at effectively zero marginal cost beyond participation incentives.

To empirically evaluate our method, we conduct 766 AI-led online interviews with US adults that lasted for close to 400 hours in total and generated 1,026,000 words of interview transcripts. The data collection required about ten days in the field, and is thus substantially faster than traditional interviews. We selected two complementary applications that account for a large share of economic research, i.e., understanding (i) economic behavior and (ii) economic policy views. Specifically, for examining the value of AI-led interviews for studying economic behavior, we interviewed non-investors about their reasons for not participating in equity markets

¹Experimental evidence suggests there are no differences in the quality and quantity of information communicated in online interviews compared to face-to-face interviews (Namey et al., 2020). An online interview might have several benefits on its own, including a greater sense of privacy and control (Gibson, 2022). Whether face-to-face conversations are needed depends on the researcher’s objectives.

²The code, instructions, and setup scripts are available here: <https://github.com/fchop/interviews>

in August 2023. For our second application, we interviewed US adults about their views on US import tariffs in May 2025. These applications provide ideal settings because of the significant heterogeneity in expected responses and the likely discrepancy in reasoning between laypeople and economists on these topics, which renders rich open-ended data especially valuable.

Interview process quality We begin our evaluation of interview quality by examining the interview process itself. As AI-led interviews are conducted without a researcher present, we cannot rely on anecdotal observation to assess what the AI interviewer is doing or how participants are responding. Furthermore, it is a common problem that LLMs often fail to follow explicit instructions (Zhou et al., 2023; Pyatkin et al., 2025), making a systematic evaluation important. To open this black box, we hand-code over 18,000 interview questions and responses for ten markers of interview quality. This task requires a human annotator to avoid circularity, as even frontier LLMs are known to have low accuracy in identifying their own mistakes (Kamoi et al., 2024). First, we find that the AI interviewer robustly adheres to key methodological principles for interviewing. For example, over 95% of questions are rated as open-ended rather than narrowly framed, perceived as neutral and non-leading rather than suggestive of any possible response, and deemed highly relevant to the interview topic. Concerns about the AI interviewer biasing responses are thus not supported by the data. Of course, interviewing requires more than simply asking open-ended questions. A skilled interviewer responds to conversational cues, clarifies ambiguous points, and follows up on notable themes. To evaluate these aspects, participants’ responses are rated on their ambiguity and need for further clarification. We find that the probability of asking a clarification question strongly increases after statements considered ambiguous by human coders. This suggests a robust contextual understanding of what types of questions are needed. Reassuringly, we also find essentially zero instances of hallucination.

Second, we use participants’ own evaluation of the interview as a direct lens into the interview quality. Participants consistently (i) rate the overall interview as very positive, (ii) perceive the conversation as feeling very natural, (iii) perceive the questions as open-ended and unbiased, (iv) indicate that the interview allows them to accurately express their own views and opinions about the topic, and (v) report feeling comfortable during the interview. As a result, a majority of participants would prefer an AI interviewer over a human interviewer and over 96% would like to participate in another AI-led interview. These evaluations are consistent with other evidence. First, when hand-coding transcripts, we detect almost no signs of discomfort or annoyance. Second, participants exert consistently high effort when answering questions, without any decrease in the amount of text produced or the time taken to answer questions over the course of the interview. Third, we replicate a weak preference for AI-led interviews in a separate survey experiment in which we randomize whether participants are offered the chance to participate in an AI-led or a human-led interview.

Hypothesis generation We next turn to the *content* of interviews and the hallmark of qualitative research: generating novel hypotheses grounded in participants’ own accounts. We examine

whether AI-led interviews can generate novel insights into economic behavior and reasoning. Across our two applications, we demonstrate the usefulness of AI-led interviews for qualitative research purposes. Specifically, we conduct an inductive thematic analysis of our interview transcripts with the goal of discovering notable themes *from the data itself* rather than forcing categories onto the data from a priori theory (Glaser and Strauss, 2017). We borrow techniques from sociology to ensure that the analytical cost for the researcher do not increase with the data, while maintaining full human control over the interpretation of the raw data (Nelson, 2020). The output of this analysis is a coding manual with a set of codes representing notable themes in the data. We hand-code transcripts based on the coding manuals and validate that this task can be delegated to an LLM at low cost by validating LLM-assigned codes against human annotators.

In both applications, we recover a wide range of systematic themes, including both novel aspects and those consistent with previous research. Perhaps most interestingly, a recurring finding across our applications is that AI-led interviews uncover novel *mental models* and economic mechanisms shaping people’s views about the world. In the tariff interviews, for example, we find that disagreement between supporters and opponents is not necessarily about whether tariffs have costs—nearly all participants acknowledge that tariffs raise prices in the short term—but about how they reason through the economic mechanisms shaping longer-term outcomes. Opponents, for example, point to structural barriers and the competitive advantages of other countries to explain why tariff-induced job creation is unlikely. Supporters, by contrast, argue that consumer substitution effects and innovation incentives will mitigate adverse economic damages and render the US economy more competitive in the long run. Our interviews thus help identify the causal reasoning underlying people’s policy views.

Similarly, in the stock market interviews, we find that many participants hold a mental model that leads them to overestimate the actual cost of participating in equity markets. Participants believe that successful investing requires *actively* selecting individual stocks, constant portfolio monitoring, and quick decision-making in response to ever-changing market prices. As a result, people mistakenly conclude that owning stocks is necessarily a time-consuming and stressful activity that is not justified by the expected returns. Our interviews thus provide a strong argument for the relevance of misperceived participation costs in driving non-participation.

We then demonstrate the external validity of this discovery out-of-sample. Specifically, we conduct a separate survey ($n = 1,000$) with investors and non-investors that includes tailored survey questions to directly measure the presence of this mental model. In support of high external validity, we find that the mental model is 22 p.p. more common among non-investors. This exercise provides a proof of principle that AI-led interviews can generate meaningful insights into economic problems that can be validated out-of-sample with standard survey techniques.

Internal and predictive validity Next, we revisit an old critique of survey methods in the context of AI-led interviews: Do people mean what they say? In other words, how much can we trust interview responses to be genuine? We address this concern in two ways. We first show that

interview responses strongly correlate with responses to validated survey instruments designed to elicit the same construct. For example, participants score half a standard deviation higher on survey measures of risk-aversion if risk-related concerns surface during interviews. Moreover, there is a high degree of internal consistency between what participants say during interviews and how they respond to conceptually related multiple-choice questions later in the survey. For example, multiple choice options in the survey are about 50 p.p. more likely to be selected if a matching theme surfaced during the interview.

We next impose a stricter test, namely whether interview responses predict future behavior. Specifically, we conduct a long-term follow-up survey with our stock market interview sample after eight months. The follow-up survey includes financially incentivized outcomes, such as willingness-to-pay measures or portfolio choices, that we find to be strongly correlated with participants' interview responses from a half a year earlier in the expected way. For example, participants who state they are unwilling to make nominal losses in the interviews are 180% more likely to demand 100% loss protection to put actual money at risk in the stock market. We are not aware of a comparable study that demonstrates the predictive validity of interview responses for future behavior. This evidence suggests that AI-led interviews capture meaningful factors that reflect how people genuinely think and behave.

Benchmarks Finally, a key methodological question for economists is whether AI-led interviews will generate richer qualitative data than open-ended survey questions—the only other method for collecting qualitative data at scale—or whether both methods yield equivalent results. This comparison is particularly relevant as these methods have an almost identical cost structure.³ To answer this question, we conduct separate benchmark surveys that include open-ended survey questions that mirror the structure and content of the interviews. We then compare AI-led interviews with the data obtained from (i) a single open-ended survey question, (ii) multiple open-ended questions, and (iii) an AI-led “interview” without any dynamically generated probing question. We then use identical methods to assign codes for each benchmark. This set of progressively stricter benchmarks allows us to decompose the benefits of AI-led interviews into (a) the gains from asking “more” open-ended questions, (b) the effects of interacting with a conversation partner, and (c) the added-value of dynamic follow-up questions.

Our first finding is that AI-led interview generate responses from participants that are thematically richer in both a qualitative and quantitative sense. Compared to asking a single open-ended question, AI-led interviews uncover about five times as many unique (and empirically common) themes, both at the participant level and in the cross-section. Our decomposition further shows that asking more questions or accounting for differences in the response mode closes at most half of the gap to an AI-led interview. In other words, most of the added-value of

³Our goal is not to replace human-led interviews, but to examine whether AI-led interviews allow economists to collect better qualitative data *at scale*. Surveys are the preferred method for measuring ex-ante known, well-defined constructs with closed-ended questions. Our focus is, however, on how AI-led interviews compare to surveys at the task of collecting *open-ended* data.

AI-led interviews can be attributed to the dynamically generated probing question. This implies that researchers using open-ended survey questions would miss quantitatively common and economically relevant themes. Importantly, the missed themes include both conceptually novel ideas—such as the *active investing* model in the context of our stock market interviews—and themes that we have validated above to be predictive of future behavior. More generally, we find that open-ended survey questions mainly capture a narrow range of concerns that are top-of-mind, struggling to reveal *why* and *how* people reason about economic problems. By contrast, we find that AI-led interviews can uncover mental models, causal narratives, and economic mechanisms that rarely surface in responses to survey questions.

Our second finding concerns the speed at which AI-led interviews achieve thematic saturation: the point where additional interviews do not reveal any new themes at the extensive margin. A bootstrap simulation shows that AI-led interviews on average saturate in small samples of $N = 25$ that are common in qualitative research studies using human-led interviews (Guest et al., 2006). This implies that the probing questions asked by the AI interviewer are efficient in the sense that they uncover novel themes as quickly as typical human-led interviews.

Related literature Our paper contributes to several strands of research. First, we develop the first scalable AI interviewer for conducting qualitative interviews and evaluate its performance across several complementary quality metrics.⁴ We develop several metrics that can be used to annotate interview transcripts to assess interviewer quality. For example, we show that the AI interviewer adheres to best practices for interviewing, asks probing questions in line with human judgment, and refrains from hallucination. Moreover, we show that participants perceive AI-led interviews as positive, with many participants preferring AI-led interviews. A growing literature has evaluated AI-led interviews along similar and complementary quality metrics. Cuevas et al. (2025) conduct AI-led interviews and find that while the interviews generate high engagement and respondent satisfaction, a human evaluation of these interviews shows a lack of richness. Several subsequent studies have compared AI-led and human-led interviews, consistently finding that AI-led interviews and human-led interviews are rated similarly on quality (Wuttke et al., 2025; Guven et al., 2025; Geiecke and Jaravel, 2025).

Second, we contribute to research on survey methods in economics (Haaland et al., 2023; de Quidt et al., 2018; Stantcheva, 2023; Fuster and Zafar, 2023; Manski, 2004) and the large literature using open-ended survey questions as a source of qualitative data (Hübsch et al., 2026; Andre et al., 2026, 2024; Stantcheva, 2021; Bursztyn et al., 2023; Galasso et al., 2024; Graeber et al., 2024a,b; Link et al., 2024; Hager et al., 2025; Bordalo et al., 2026; Chopra et al., 2025, see Haaland et al. (2025) for a review). Our main contribution is to demonstrate and quantify that AI-led interviews generate thematically richer responses compared to asking open-ended survey questions. We show that the benefits of AI-led interviews stem from the use of probing

⁴We posted the first version of our paper online on September 15, 2023 (<https://x.com/FelixChopra/status/1702791162753696091>), which predates all other papers on AI-led qualitative interviews.

questions. Consequently, many relevant explanations for economic behavior or political views would go unnoticed unless AI-led interviews are used. Our second contribution is to show that these themes matter: interview responses are internally valid and predict future behavior even half a year later.

Third, we add to the qualitative literature on hypothesis generation in economics (Bewley, 1995; Shiller, 1997; Blinder et al., 1998; Bewley, 1999; Morduch and Schneider, 2017; Sivaramakrishnan et al., 2017; Duraj et al., 2025; Mutz, 2021; Bustos et al., 2022; Bergman et al., 2024; Ludwig and Mullainathan, 2024, see Starr (2014) for a review). Our contribution is to provide the first empirical evidence that AI-led interviews can be used to generate novel hypotheses about economic problems. We further contribute by showing that AI-led interviews reach saturation at a level and speed comparable to human-led interviews. Recent work uses AI-led interviews to generate novel hypotheses (Baader et al., 2025; Hwang et al., 2025; Puente, 2026)

Finally, we relate to research using LLM-powered chatbots to deliver treatment interventions at scale (Costello et al., 2024; Boissin et al., 2025; Ash et al., 2026). Jabarian and Henkel (2025) introduce voice-based AI-led interviews and find that they outperform interviews conducted by professional recruiters. Chopra et al. (2026) use our AI interviewer setup to reduce social media use by increasing people’s motivation for change. More broadly, we relate to work examining the impact of AI on economic outcomes (Noy and Zhang, 2023; Brynjolfsson et al., 2025) and economic research (Horton, 2023; Korinek, 2023; Charness et al., 2025; Ash et al., 2024).

The remainder of the paper is organized as follows. Section 2 provides details about the AI interviewer and our empirical applications. Section 3 examines the interview process before Section 4 presents our evidence that AI-led interviews can generate novel hypotheses. Section 5 examines the internal and predictive validity of interview responses and Section 6 benchmarks AI-led interviews against other survey methods. Finally, Section 7 concludes.

2 Methods and applications

We explain the scalable implementation of AI-conducted interviews in Section 2.1 and describe the design and interview structure of our empirical applications in Section 2.2. Methods for large-scale text analysis of these interview transcripts are presented in Section 2.3.

2.1 Conducting interviews at scale with an AI interviewer

Qualitative interviews differ from surveys because participants can express their views in their own words in an unconstrained way, and interviewers have the opportunity to ask *adaptive* follow-up questions—in short: flexible, open-ended interaction. This flexibility offers three advantages over surveys, addressing common limitations of stand-alone open-ended survey questions. First, while responses to stand-alone open-ended survey questions tend to be short and ambiguous, interviews generate detailed conversational records that offer the necessary context

for interpretation and classification. Second, interviews comprehensively measure how people reason and interpret the world, and thus do not limit the researcher to participants’ immediate, top-of-mind reactions. Third, interviews facilitate the discovery of novel hypotheses as unanticipated themes can emerge since—in contrast to surveys—questions are less predetermined by the researcher’s prior assumptions about what is relevant or important.⁵


Design objectives Our goal is to maintain the above benefits of interviews when conducting interviews at scale with an *AI interviewer*. In this paper, we focus on *semi-structured interviews* (Knott et al., 2022), where researchers define an interview topic guide in advance but retain the freedom to adjust their line of questioning within each topic based on interviewees’ responses. This structure ensures comparability across interviews while preserving the interview’s exploratory strengths (DeLuca et al., 2023). To maintain these benefits at scale, an AI interviewer must meet three core design objectives. First, it must be capable of conducting high-quality interviews—asking neutral, open-ended questions, but also engaging in adaptive probing that responds meaningfully to what the participant says. This requires reading between the lines, responding to contextual cues, revisiting relevant earlier points, and introducing new topics without disrupting the flow of the conversation. Second, it must be practically reliable: robust to unusual input, while fast and consistent in generating high-quality output. Finally, from an ethics perspective, the AI system must be safe for use with human participants. We next describe how we meet these objectives in practice.

Modular system: Single or multi-agent workflow Our AI interviewer is surprisingly simple and modular, consisting of a database, a set of prompts, and a large language model. At any stage of the interview, the next question is dynamically generated by *prompting* an LLM via an Application Programming Interface (API). We started with OpenAI’s GPT-4 base model in 2023 (without further fine-tuning the model), and switched to GPT-4o in 2025, but other frontier or open-source models can be used if preferred.⁶ The workflow for generating the next question is shown in Panel A of Figure 1. Researchers can switch between a (simpler) single-agent or a multi-agent workflow with more fine-grained control by activating agents indicated as optional. After participants answer a question, the survey software forwards the response to the AI interviewer (a Python application hosted on a server). To generate the next question, up to four distinct tasks are carried out by up to five different *agents*: the probing agent (required), the moderator agent (optional), the summary agent (optional), and the topic agent (optional). Each agent is an LLM with a unique *prompt* of about 150-350 words, i.e.,

⁵As a detailed review and discussion of interviews and qualitative methods is beyond the scope of this paper, we defer the interested reader to the many excellent textbooks and overview articles that have been written on these subjects (e.g., Creswell, 2009; Mohajan, 2018; Small and Calarco, 2022; Knott et al., 2022).

⁶Frontier LLMs can conduct interviews in other languages beyond English, provided the language is well-represented in the training data. For example, our system has been used to conduct interviews in German with high respondent satisfaction (Baader et al., 2025). Moreover, they can understand and respond to text with severe spelling errors and non-standard grammar, making them particularly useful for interviewing participants with lower literacy levels.

Figure 1: The AI interviewer and the participant experience



Note: Panel A provides an overview of the workflow of how the AI interviewer generates the next interview question. Agents marked optional can be individually deactivated to turn the multi-agent workflow into a single-agent workflow. In this case, the probing agent handles both topic transitions and the generation of follow-up questions. Panel B provides a screenshot of the 2025 chat interface with the option to either answer in writing or by recording a voice message that is transcribed automatically.

short written instructions describing the desired behavior and output. Together, they collectively function as our AI interviewer. This *decomposed prompting* approach that splits the complex task of interviewing into simpler tasks for different agents is a common strategy to improve LLM performance (Khot et al., 2023). Our modular approach allows researchers to add or subtract agents as needed. For instance, a fully unstructured interview without a topic guide can be accommodated by deactivating the topic agent. For a longer interview, the summary agent can be activated to compile analytical notes on the fly as input for other agents.

The workflow for generating questions starts with the **moderator agent (optional)** that verifies whether the participant’s message includes harmful content, sensitive personal information (e.g., names, credit card details), or attempts to steer the conversation off topic. Flagged messages are not stored and participants are gently nudged to rephrase or decline to answer. The interview stops if messages are flagged too often.⁷ After the moderator agent, the next step is to consult the interview’s **topic guide**. Researchers can fully customize the interview by defining a set of topics

⁷In practice, fewer than 1% of messages are flagged, and most flagged messages are false positives. Note that concerns about participants disclosing sensitive information, although instructed not to, apply equally to open-ended survey questions that lack the safeguard described above. OpenAI’s API neither stores nor uses data for training.

that are to be covered sequentially with a question “budget” per topic, thus giving researchers full control over their length (see Table A.1 for examples). Once this question budget for the current topic is exhausted, the conversation moves to the next interview topic.⁸ When transitioning to a new topic, the **summary agent (optional)** starts by updating an analytic summary of the conversation so far for the other agents. The summary agent is designed for very long interviews that require compressing the text input to focus the model’s attention. After the summary is updated, the **topic agent (optional)** is responsible for introducing the next interview topic. The agent is instructed to find a smooth and natural transition to bridge what has been discussed with what will be covered next. If no topic transition is needed, the next question is instead generated by the **probing agent** (the only required agent in our setup). The probing agent receives the current interview topic and the previous conversation (or, optionally, the analytic summary) as context for thinking about a probing, clarification, or follow-up question to achieve both depth and breadth of coverage. In a final step, the **moderator agent (optional)** reviews and, if needed, stops potentially problematic questions from being asked as a final safeguard.

Prompts As qualitative methods textbooks are part of their training dataset, frontier large language models already have a strong baseline understanding of qualitative interviews, and thus require only minimal instructions and reminders to activate this context. For each agent, we curate individual prompt templates of no more than 350 words. These prompt templates include best practices for interviewing that we compiled from standard textbooks (Legard et al., 2003; Rubin and Rubin, 2011; Knott et al., 2022; Corbin and Strauss, 2014). The prompts for each agent follow the same basic structure. We first introduce a context where the AI interviewer is informed that it is an “AI proficient in conducting qualitative interviews for academic research” and inform the AI about the general topic and context of the interview. Establishing a role for an LLM (“you are an. . .”) has been shown to increase the performance of LLMs (Kong et al., 2024). After providing the previous conversation, analytic summary and information about the current topic as inputs, we describe the task and guidelines. To make this concrete, we present the core part of the probing agent’s prompt below and relegate the other prompts to our GitHub repository (see Section D.1):

TASK: Your task is to formulate the next probing question for the Current Conversation. The question should align with the Current Interview Topic, helping us to better understand *[insert interview topic]*.

GENERAL GUIDELINES:

1. Open-endedness: Always craft open-ended questions ("how", "what", "why") that allow detailed and authentic responses without limiting the interviewee to "yes" or "no" answers.
2. Neutrality: Use questions that are unbiased and don’t lead the interviewee towards a particular answer. Don’t judge or comment on what was said.

⁸When deactivating the topic agent, the probing agent endogenously determines the number of questions per interview topic depending on how the conversation unfolds. However, we argue that a *deterministic* workflow might often be preferred: Fixing the number of questions renders interviews more comparable and offers more *control*, avoiding the risk of having some very short or very long interviews. This is practically important when participant recruitment requires announcing the expected study duration in minutes ex-ante.

3. Respect: Approach sensitive and personal topics with care. If the interviewee signals discomfort, respect their boundaries and move on.
4. Relevance: Prioritize themes central to *[insert interview topic]*. Don't ask for overly specific examples, details, or experiences that are unlikely to reveal new insights.
5. Focus: Generally, avoid recaps. However, if revisiting earlier points, provide a concise reference for context. Ensure your probing question targets only one theme or aspect.

PROBING GUIDELINES:

1. Depth: Initial responses are often at a "surface" level (brief, generic, or lacking personal reflection). Follow up on promising themes hinting at depth and alignment with the research objective, exploring the interviewee's reasons, motivations, opinions, and beliefs.
2. Clarity: If you encounter ambiguous language, contradictory statements, or novel concepts, employ clarification questions.
3. Flexibility: Follow the interviewee's lead, but gently redirect if needed. Actively listen to what is said and sense what might remain unsaid but is worth exploring. Explore nuances when they emerge; if responses are repetitive or remain on the surface, pivot to areas not yet covered in depth.

As qualitative disciplines and scholars strongly differ in their views on how to “best” conduct interviews, the above prompt template focuses on very minimal aspects that most researchers would agree on (e.g., open-ended, non-leading questions, clarifying ambiguities, refraining from judgment, and so on). If the AI interviewer should follow a particular style of interviewing, researchers can modify our prompt templates in that direction. The above prompt contains no application-specific instructions, making it portable across domains.⁹

Interview interface: Text or voice input Participants interact with the AI interviewer through a chat interface (Panel B of Figure 1) that can be embedded into standard survey software, making it compatible with both online interviews and computer-assisted personal interviews in the field. It resembles common text messaging applications to minimize technological barriers for participants. Participants can respond to questions either by typing directly into the text box or by recording a voice message that is transcribed with OpenAI's Whisper API (without storing the audio message). Participants can thus choose the input method—audio or text—that best allows them to convey their thoughts.¹⁰ As of 2025, the waiting time until the next question is displayed is very short and essentially negligible with a median duration of 1.3 seconds.

Open-source software To facilitate adoption by other researchers, our AI interviewer application is open-source with detailed documentation and setup scripts to run the model either locally or hosted on a server.¹¹ The local setup is ideal for experimentation and adaptation of the setup to a new research context. For eventual data collections, we provide production-ready code that can be deployed on standard cloud infrastructures (e.g., Amazon AWS for about \$1 per month). We provide custom code to integrate the chat interface into standard Qualtrics surveys (by copy-pasting HTML and JavaScript into a “Text/Graphic” question). Integrating the AI-led

⁹The total LLM usage costs for designing the AI interviewer and developing our prompts in 2023 were negligible, amounting to less than \$50 in 2024 prices.

¹⁰Evidence from randomized controlled trials suggests that spoken responses yield content of comparable richness to written responses (Namey et al., 2020), a finding that we replicate in our second application where the voice input feature was added to our setup (see Appendix Section C.4).

¹¹The full code and documentation are available here: <https://github.com/fchop/interviews>

interview directly into the survey environment reduces respondent attrition by eliminating the need to navigate away from the survey and increases flexibility by allowing customization of prompts based on previous survey responses.

2.2 Applications: Design and interview structure

The remainder of the paper focuses on a comprehensive examination of AI-led interviews, where we evaluate the following four key questions. First, do AI-led interviews keep respondents engaged and satisfied throughout the interview? Second, does the AI interviewer adhere to its instructions, securing alignment with best-practice interview standards? Third, do the interviews produce high-quality data that provide rich insights into economic behavior and economic reasoning? And fourth, do AI-led interviews add substantial value beyond surveys with open-ended questions? To investigate these questions, we conduct about 800 AI-led interviews in two complementary domains for understanding (i) individual economic behavior and (ii) economic policy views. Together, these applications cover a large share of experimental and survey-based research in economics. We describe the study design and interview structure of our concrete applications below, and introduce additional data collections throughout the paper when relevant. Table A.2 provides an overview of all studies and Section F contains the experimental instructions.

2.2.1 Economic behavior: Stock market non-participation

To investigate whether AI-led interviews generate valuable data to study economic behavior, we selected the stock market non-participation puzzle as our first application and ask non-investors why they do not participate in equity markets (Guiso and Sodini, 2013). This puzzle is of great interest to economists and policymakers, and the substantial literature on the puzzle sets a high bar for AI-led interviews to potentially generate a novel hypothesis about the underlying drivers.

Sample Our analysis sample consists of 381 interviews with US adults conducted between August 23 and September 5, 2023.¹² Participants were recruited from Prolific, a survey platform commonly used in economic research that is associated with high data quality (Peer et al., 2022). We use a screener survey with standard requirements to recruit the relevant subpopulation of non-investors with the financial means to participate in equity markets, focusing on those without any direct or indirect stock holdings (see Section F.1.1 for details). To ensure a high quality of responses as required for interviews (Rubin and Rubin, 2011), the screener survey includes a standard attention check and additionally screens out participants who demonstrate an unwillingness to engage with an open-ended survey question. We exclude 14 participants flagged during a human review of all interview transcripts as having used AI tools. Table A.3 provides summary statistics for our sample, which deviates from the general population as we target

¹²The field time is thus substantially shorter compared to sequentially conducted interviews with a human interviewer, and only required several days due to OpenAI’s API rate limit in 2023.

non-investors. Our sample is close to the US population in terms of age, but has a higher share of women and participants with a college degree, as is common for online panels (Armantier et al., 2017). The average household income is \$72,000, with 60% reporting liquid savings equivalent to at least two months of income. Reassuringly, net financial assets are about \$51,000, indicating that most participants have the means to invest in stocks.

Design and interview structure All interviews were conducted in writing (without our voice input feature) as part of an online survey with a median completion time of 40 minutes and a fixed compensation of \$12. The survey starts with a detailed consent form disclosing that the study includes a chat-based interview with an AI chatbot to allow participants to select out of the study if they feel uncomfortable with either an interview or engaging with a machine (a selection decision we will study in Section 3). We also inform them that their responses will be processed via OpenAI’s API, but that their data will neither be used for model training nor permanently stored by OpenAI. Before the interview, we elicit background characteristics using standard survey questions and allow participants to familiarize themselves with the chat interface. After the interview, we elicit detailed feedback on the interview experience and include additional survey questions on factors correlated with stock market non-participation.

The interview itself is semi-structured with four conversational topics, with the verbatim instructions provided to the AI interviewer for each topic shown in Table A.1. For this data collection, the AI interviewer consisted of a summary agent, a topic agent, and a probing agent.¹³ The conversation always starts with a broad question that signals interest in the participant’s personal reasons for non-participation: “I am interested in learning more about why you currently do not own any stocks or stock mutual funds. Can you help me understand the main factors or reasons why you are not participating in the stock market?”¹⁴ After five follow-up questions, the conversation moves to the next topic that instructs the AI interviewer to explore the “perceived barriers or challenges preventing [the respondent] from participating in the stock market.” We chose this topic to ensure that the interview covers other relevant reasons for non-participation, as it naturally “reorients” a conversation stuck in a single line of inquiry. The third topic asks participants to walk the interviewer through a hypothetical scenario in which they own stocks (three questions), employing the commonly used strategy of counterfactual reasoning to uncover participants’ reasoning about what it means to own stocks. The fourth and final topic revolves around “conditions or changes needed for the interviewee to consider investing in the stock market.” Following common practices (Rubin and Rubin, 2011), the interview concludes with two questions on whether participants would like to share anything else and asking them to

¹³At the time of the data collection, the relatively short context window of GPT-4 made a summary agent strictly necessary. While current frontier models have longer context windows, a summary agent might still be useful to guide the model’s attention in very long interviews. Furthermore, the final conversation summary can be shown to the interviewee, allowing them to rate whether the interview accurately captured their views.

¹⁴While interviews usually start by building rapport, we deliberately chose this question because it resembles the type of open-ended survey questions economists use to gather qualitative data. We later contrast the themes revealed in participants’ responses to this question with the themes a researcher would uncover based on a full interview.

reflect on the conversation and again identify their main reason for non-participation. After the interview, we present the conversation summary generated by the summary agent and ask participants whether it accurately represents their views expressed in the interview.

2.2.2 Economic policy views: US import tariffs

To examine the value of AI-led interviews for studying economic policy views, we interview US adults about their views on the US import tariffs imposed by President Trump in 2025. Tariffs are a policy issue on which laypeople’s accounts are likely to diverge from standard economic theory, rendering open-ended qualitative data especially valuable (Mutz, 2021).

Sample In May and June 2025, we conducted 385 online interviews with US adults recruited from Prolific with a screener survey that elicits baseline support for tariffs and includes similar quality requirements as described above.¹⁵ 30% of our sample support and 70% oppose the 2025 US import tariffs.

Design and interview structure Interviews are again conducted as part of an online survey. We measure background characteristics before the interview and participant feedback, as well as an additional survey module eliciting views about tariffs using structured questions, after the interview. For this data collection, we used the moderator, probing, and topic agents. Moreover, a key difference from our first application is that participants could now answer interview questions either by typing their response or by recording a voice message that we automatically transcribe. The interview itself is again semi-structured, following the topic guide outlined in Table A.1 with the opening question: “First of all, what are the main reasons you support or oppose the new 10 percent tariff on imports?” After five follow-up questions, the AI interviewer uses a counter-perspective prompt, asking whether there are any positive (negative) aspects of tariffs if the participant previously discussed primarily negative (positive) aspects. The third part asks participants to explain to someone unfamiliar with tariffs how they work, while the fourth part explores participants’ beliefs about the consequences in a scenario where the US raises tariffs to 30%. Finally, the AI interviewer is tasked with reviewing the conversation and following up on promising but unexplored aspects before concluding the interview.

2.3 Methods for analyzing interviews

This subsection explains the general methods we use to (1) qualitatively analyze transcripts and inductively construct a coding manual, (2) conduct human coding of interviews, and (3) validate

¹⁵We take comprehensive measures to mitigate growing concerns about bots and LLM use in online surveys. To prevent the use of AI tools, we programmatically block participants from copying or pasting text, rendering LLM tools like ChatGPT ineffective, since interview question and LLM responses would have to be manually transcribed between the survey interface and the tool, raising the cost of tool use above the level of answering the interview question directly. In addition to Captcha/reCaptcha methods, we track—among others—the ratio of text produced to unique keystroke events to identify bots in our screener survey that “produce” text without hitting a physical keyboard. A final human review of the interview transcripts did not find any evidence of remaining bot activity.

AI coding of interview transcripts against a human benchmark. All methods are designed to scale to a text corpus of millions of words while retaining a “human in the loop.” Similar to how applied researchers are expected to understand their raw data and setting, it seems prudent for those working with interviews to stay close to the source text to avoid losing nuance and context during interpretative aggregation (see Nelson, 2020, for the full argument).

2.3.1 Inductive thematic analysis

We use inductive thematic analysis to analyze our transcripts, i.e., we aim to discover notable themes solely from what participants say rather than imposing what we as researchers believe should be present in the data (Knott et al., 2022). This approach is deliberately agnostic and aids the discovery of novel ideas in contrast to deductive methods that derive codes from theory (Glaser and Strauss, 2017). A key output of this process is a collection of *codes* that each represent a notable, recurring theme (e.g., concerns about the costs of owning stocks, or concerns about tariff-induced inflation). The codes jointly define a so-called *coding manual*, which can be used for labeling responses. To arrive at a scalable method, we borrow techniques from computational sociology (Nelson, 2020). First, we independently read a random sample of 50 interview transcripts to obtain a robust baseline understanding of the context and the language participants use. In a second step, we transform interview transcripts into a text corpus where each document is a question-answer pair from an interview (about 13,000 pairs per application). We then use BERTopic, an unsupervised embedding-based clustering algorithm, to sort the data into clusters of semantically similar text (Grootendorst, 2022). The motivation for this step is that it leverages the pattern recognition ability of machine learning to increase the chance and speed of discovering a notable theme in the data (Ludwig and Mullainathan, 2024). In contrast to bag-of-words or word list methods, BERTopic constructs more semantically coherent topic clusters than non-embedding-based methods that ignore the semantic relationship of words. In a third step, we independently and sequentially review each cluster, focusing on the model’s top-ranked responses. For notable themes, we write a short synopsis, code definition, and compile representative example quotes. At the end, we compare notes, consolidate our list of codes, and ensure a shared conceptual understanding of codes. The above process, while described linearly, was iterative in practice.¹⁶

2.3.2 Coding interviews

Human coding As a ground truth for validating the quality of AI coding, a team of research assistants trained in qualitative interview methods hand-coded over 18,000 questions and responses from our interview transcripts. Appendix Section E provides details about the training, calibration of annotators, and coding practices. The interrater agreement between research assistants is 83%, indicating a high quality of the hand-coded data. We annotate interviews twice,

¹⁶By *not* using LLMs to generate the coding manual, we mitigate concerns about LLM biases affecting the analysis. In Section 2.3.2, we validate that these human-discovered themes can be accurately classified with LLMs.


once using the coding manual derived from the inductive thematic analysis, and another time using a second coding manual focused on rating interviews based on explicit quality criteria for interviewing. We fully hand-coded all stock market interviews, while a random half of all tariff interviews are hand-coded for the analyses presented in Section 3.

AI coding and validation Hand-coding is the gold standard for annotating qualitative text data, but the time and labor costs render it ineffective for working with interviews at scale. To address this, we demonstrate that researchers can use LLMs to automate the hand-coding process. Of course, it remains crucial to validate the performance of AI coding against a human benchmark to guard against hallucination and systematic model biases (Carlson and Dell, 2025). For AI coding, we use OpenAI’s GPT-4o model. Section E.2 contains the prompt, which includes general coding instructions, such as avoiding implicit assumptions, a placeholder for the definition of the code, and a list of example responses. To increase accuracy, we ask the model to supplement its decision with a justification and to cite evidence from the transcript. Using data from the stock market interviews, we find that AI coding closely aligns with the human-coded ground truth, as shown in Figure B.1. The global interrater agreement between AI and human coders across all codes is 86%, and thus as high as the human–human agreement in our setting. At the code level, the agreement ranges from 81% at the 25th percentile to 92% at the 75th percentile, with GPT-4o on average being 6.4 p.p. more likely to assign a code to an interview than a human. Taken together, this suggests that human-validated AI coding can serve as a cost-effective substitute for labeling interview transcripts.

3 Behavioral and perceptual markers of interview quality

Before turning to the thematic content of our interviews, we first examine the quality of the interview process itself. This step is essential: because AI-led interviews are conducted without direct researcher involvement, we cannot rely on anecdotal observation to assess what the AI interviewer is doing or how participants are responding. Instead, we assemble a set of behavioral and perceptual markers that should be indicative of high interview quality. These include the time and effort participants invest (Section 3.1), the quality of the questions asked by the AI interviewer as judged by trained human annotators (Section 3.2), participants’ subjective evaluations of the interview (Section 3.3), and an experiment measuring demand for AI-led interviews (Section 3.4). What unites these exercises is that they do not depend on analyzing the content of participants’ responses. Instead, they offer an external lens on the interaction itself. These process-level metrics are a necessary foundation before turning to the content-based analyses in Section 4, 6, and 5. They speak to the quality of the interview as an interaction, i.e., whether it fosters engagement, adheres to methodological best practices, and creates the conditions under which meaningful qualitative data can be obtained.

Figure 2: Text production and thinking time are high and stable over the course of the interview



Note: Data from the stock market and tariff interview studies. Panel A: Histogram of the interview duration in 5-minute bins and censored at 80 minutes. Panel B: Average time to answer the k^{th} question in the interview in seconds. For the stock market interviews, the response time to the first question could not be recorded yet. Panel C: Histogram of the total text written by participants in thousands of characters and censored at 8,000 characters. For the tariff interviews, the mass at 8,000+ characters is driven by participants who used the voice-to-text input method. Panel D: Average text written in response to the k^{th} question in the interview in characters. Dashed lines in Panels C and D indicate that questions 17 and 18 are the closing questions of the interviews.

3.1 Interview duration and participants' effort

We first present descriptive evidence that participants take the interviews seriously and exert high effort answering the questions. As shown in Panel A of Figure 2, many participants spend considerable time completing the interview. The stock market interviews take about 33 minutes on average to complete, and participants produce 654 words (3,487 characters) of text during this time (Panel C). With 29 words per minute, the text output is comparable to the benchmark of about 20 words per minute obtained in chat-based interviews with human interviewers (Namey et al., 2020). The tariff interviews are slightly shorter with a median duration of about 25 minutes, but text production exhibits a significant right tail, which is partly driven by participants who used the voice-to-text input feature (see Appendix Section C.3 and C.4 for further results showing that interviews increase text production on a question-by-question basis).

Over 90% of participants who start the survey complete the full interview, which is striking given that participation in the interview was voluntary and participants knew payments were not conditional on participation. In fact, respondents could decide to proceed in the survey and skip the interview after a five-minute delay. The very low attrition rate is encouraging given concerns that survey respondents might not exert effort without meaningful incentives, especially for open-ended questions that require more effort from participants compared to closed-ended questions (Haaland et al., 2023). Notably, effort as measured by text output or thinking time does not decline over the course of the interview. Participants consistently take about 100–120 seconds to answer interview questions (Panel B), writing about 200–300 characters on each

Table 1: Characterizing the behavior of the AI interviewer

	(1) Economic behavior (Stock market)	(2) Policy views (Tariffs)
<i>Local properties</i>		
Open-ended question asked	94.8%	99.8%
Question relevant to the interview objective	95.0%	97.9%
Neutral and non-leading question	95.7%	99.6%
<i>Contextual properties</i>		
Novel, non-repeated questions	82.7%	92.2%
Follow-up about the participant’s last answer	77.6%	37.5%
Follow-up about something mentioned earlier	15.4%	13.9%
<i>Ethics and safety</i>		
Signs of high conversational engagement	79.6%	85.2%
Annoyance with the interviewer	2.5%	0.7%
Any potential signs of discomfort	0.2%	0.1%
LLM hallucinated	0.1%	0.4%

Note: Data from the stock market and the tariff interview studies (see Table A.2). Summary statistics derived from hand-coding by a team of trained research assistants. All percentages are computed at the question-answer pair level, excluding the predetermined closing questions. Local properties of questions are assigned by judging the question in isolation, while contextual properties require knowledge of the previous conversation.

question, with the only understandable drop for the closing questions that signal the end of the interview (Panel D).

3.2 Assessing the behavior of the AI interviewer

We next open the black box of what happens during interviews and provide descriptives to answer the relevant question: What is the AI interviewer doing? To do so, we follow best practices and use human judgments as the gold standard for evaluation (Dell, 2025).¹⁷ As described in Section 2.3.2, we systematically hand-coded over 8,000 question-answer pairs on 10 markers of interview quality. These markers capture both local properties of questions (e.g., using open-ended rather than closed-ended wording) and contextual properties that depend on the previous conversation (e.g., whether a question is novel or rephrases an earlier question). Details about the coding manual are shown in Appendix Section E.

Table 1 presents the results. Reassuringly, we find that the AI interviewer consistently adheres to core methodological guidelines for interviewing across both applications. First, about 95% of questions in the stock market interviews are coded as being (i) open-ended rather than closed-ended, (ii) relevant to the interview’s objective, and (iii) non-leading, indicating that the AI does not ask biased questions. For the tariff interviews—which were conducted with a more

¹⁷It is a common problem with large language models that they often fail to follow explicit instructions (Zhou et al., 2023; Pyatkin et al., 2025). Furthermore, Kamoj et al. (2024) find that even frontier models such as GPT-4 detect errors “at very low recall” and are “much worse than humans” at error detection, making it essential to use trained *human* evaluators to evaluate the quality of the interviews and LLM alignment with instructions.

advanced LLM—these shares are close to 100%, confirming our expectations that interview quality will naturally improve over time as foundational models naturally mature.

Second, turning to contextual properties, we find that 83% (92%) of stock market (tariff) questions are unique and thus not a rephrasing of an earlier question, which research assistants were instructed to conservatively classify as “redundant”—although repetition can be useful for revisiting earlier points, clarifying the meaning of the question, or prompting a participant to tell us more. These markers of question quality capture core principles for interviewing that we *explicitly* described in our prompts for the AI interviewer, and thus speak to the AI interviewer’s instructional alignment.

An important skill for conducting interviews is understanding contextual cues and establishing a *dialogue* where questions dynamically respond to participants’ answers in a meaningful way. For the stock market interviews, we find that 76% of all questions explicitly refer to something that the participant said in their most recent response; for example, by asking a follow-up question about a newly raised point. This share is lower for the tariff interviews, where we made a minor prompt change to reduce the tendency to paraphrase the participant to make questions shorter. Moreover, 14–15% of interview questions revisit a key theme raised by the participant during earlier parts of the conversation.


Third, we assess markers of participants’ well-being and safety, arguably key objectives for interviewing. Research assistants hand-code responses for any potential signs of discomfort or annoyance with the interviewer. Reassuringly, signs of discomfort are extremely infrequent. Only 2.5% of responses indicate annoyance in our stock market interviews. These instances were typically due to respondents feeling they had already answered a question, which explains why the share is only 0.7% for the tariff interviews that include fewer repeated questions. Similarly, we find that the share of responses rated as displaying signs of high conversational engagement—which focused on whether participants make an effort to actively participate in the conversation rather than the amount of text produced—ranges from 80% to 85% across applications.

We also address concerns about model hallucination (Zhang et al., 2023), which cannot be credibly addressed without a human coder carefully analyzing each interview transcript. We detect only five (seven) instances of questions that include incorrect references to prior responses in our stock market (tariff) interviews, suggesting that hallucination is effectively absent in our contexts.¹⁸

Finally, we provide evidence that the AI interviewer also displays an understanding of when clarification questions are needed. The probability of asking a clarification question increases sharply when interviewees’ responses are hand-coded as requiring further clarification. Specifically, clarification questions are rare (2.8%) when humans perceive no need for further

¹⁸We coded this conservatively. Most cases are minor hallucinations where the interviewer made an inference error based on the previous conversation, such as concluding that the participant is interested in learning more about the stock market when this was not explicitly said.

Figure 3: Participant evaluations and perceptions of the AI-led interview



Note: Panel A uses data from the stock market interviews study ($n = 381$), while Panel B uses data from the tariff interview study ($n = 385$; see Table A.2). “Overall rating” shows participants’ overall evaluation of the interview. “Naturalness of interview” is participants’ assessment of the naturalness of the conversation. “Interviewer preference” is the response to the question whether they would prefer a human or an AI interviewer in a future interview. Note that the “neutral” response option was removed for the tariff interview study. “Accurate representation” is the share of participants approving the summary of the interview as accurate in Panel A, and shows the share saying that the interview somewhat or very accurately captures how they think about tariffs in Panel B. “Participate again” is the share that would like to participate in another AI-led interview.

clarification, but the share increases to 31% and 51%, respectively, when clarification is deemed “somewhat relevant” and “very relevant.” This result suggests that the AI interviewer is able to recognize ambiguous responses and follow up with a clarification question.

Taken together, these findings provide a direct lens into AI interviewer behavior, indicating that the AI interviewer consistently follows key methodological instructions for interviewing and displays high contextual awareness during conversations, while ensuring participant safety.

3.3 Participant perceptions of the interview

After the interview, we elicited participants’ perceptions and evaluations of the interview to obtain direct evidence on whether the interviews are perceived as a positive and natural experience by the interviewees. Feedback was elicited outside the chat interface as part of the remaining survey to facilitate honest evaluations (Nass et al., 1994). Focusing first on the overall assessment, Figure 3 shows that participants in both applications provide overwhelmingly positive reviews, with about 60% rating the interview as “very” or “extremely positive” across contexts. When

asked about the naturalness of the conversation, most participants rate the conversation as feeling natural. We also asked participants whether they would prefer a human or an AI interviewer in a future interview on a comparable topic. For our stock market interviews, 53% indicated a preference for an AI interviewer compared to only 21% who prefer a human interviewer. For the tariff interviews, we removed the neutral option, which increases the share of participants who prefer an AI interviewer to 71%. This suggests that the *experience* of participating in an AI-led interview must be sufficiently positive. Indeed, almost all of our participants (96–97%) indicate that they would like to participate in another AI-led interview. Note that the results from the 2025 tariff interviews show that the positive evaluations are not driven by the initial novelty of LLM chatbots in 2023. Table A.5 further shows that the positive evaluations of the AI interviewer and the preference for AI-led interviews is robust across sociodemographic groups defined by age, gender and education.

Finally, we show that AI-led interviews are able to accurately capture participants’ perspectives. After the stock market interviews, we asked participants to approve a summary of why they are not participating in the stock market: 99% approve the summary as accurately capturing their case. After the tariff interviews, we ask participants to rate whether the interview “accurately captures how [they] personally think about the topic of tariffs” on a 5-point Likert scale and obtain similar results. Figure B.2 provides additional direct evidence on how participants perceive the interview questions in the tariff interview, showing that questions are rated as (i) open-ended, (ii) neutral, and that participants (iii) express feeling comfortable during the interview.

3.4 Selection experiment: Revealed demand for AI-led interviews

Design The above findings suggest that people display a high willingness to engage with an AI interviewer and might even prefer AI interviews over human interviewers. To mitigate potential concerns about social desirability bias or positive selection into our AI interview samples, we ran a correspondence study with 500 US adults recruited from Prolific in November 2023 using equivalent recruitment criteria (see Section F.1.5 for instructions). Table A.3 provides summary statistics for the sample. In the survey, we inform participants that they might be eligible to participate in three additional studies. We describe each study on a separate survey page in randomized order. Each study is described analogously to how we described the stock market interview study on Prolific. However, we exogenously vary whether the study is described as (i) “a 40-minute survey,” (ii) “a 40-minute study with a text-based interview with an AI,” or (iii) “a 40-minute study with a text-based interview with a human.” On the page where the study is described, participants indicate whether they would like to be invited to the study if we decide to run it.

Results As a clean comparison without spillover, we focus on the first study described to participants. The demand for study participation is identical for a regular survey (89%) and a survey with an AI-led interview (90%). If anything, demand is 6.6 p.p. *lower* for a survey with

a human-led interview ($p = 0.078$). Moreover, exploiting within-subject variation, we find no evidence of differential selection into surveys compared to AI interviews (as shown in Table A.4). We thus replicate a weak preference for AI-led compared to human-led interviews also from an ex ante perspective prior to participation in an AI-led interview.

4 Qualitative findings and hypothesis generation

Having established that AI-led interviews maintain respondent effort and satisfaction while also adhering to best-practice interview standards, we now turn to whether they also produce rich, high-quality interview data that can be used for hypothesis generation—the hallmark of qualitative research. We present our main *qualitative* findings from the stock market interviews in Section 4.1 and then present an out-of-sample validation of our derived hypothesis in Section 4.2. In Section 4.3, we present key findings from the import tariff interviews. We return to these findings in Section 6 to show that many relevant themes would not have been present in responses to standard open-ended survey questions.

As our primary goal is to evaluate AI-led interviews as a new method for data collection, the qualitative results below are not intended as a complete account of all themes present in the interviews. Rather than focusing on the most frequently mentioned themes, we highlight those that reveal interesting perspectives relative to previous work or that illustrate the added value of AI-led interviews over open-ended survey questions (see Section 6). In each application, we present findings narratively and illustrate them with representative quotes drawn from participant responses. To remain close to the conventions of qualitative research, we follow Bergman et al. (2024) and focus on a rich description of the key mechanism with illustrative quotes from the interviewees. We refer the reader to the relevant figures in each section for code frequencies.

4.1 Main qualitative findings from the stock market interviews


Our interviews uncover considerable heterogeneity in the perceived barriers to stock market participation, confirming that non-participation is a multi-layered and multidimensional phenomenon. Our evidence supports leading explanations of non-participation, such as low financial literacy, risk-related concerns such as fear of losses, peer effects, participation costs, or income constraints (Guiso and Sodini, 2013). As the focus of our paper is methodological, we focus on three main qualitative patterns. Figure 4 presents code frequencies.

4.1.1 Mechanism 1: Participation costs and the mental model of active investing

When asked about their reasons for not participating in the stock market, a recurring theme is that respondents describe concerns about the time commitments and the work involved in owning stocks. As one respondent puts it:

Managing investments in the stock market feels like it could almost be a part-time job given the level of attention it demands. My current job is already time-consuming and mentally exhausting.

Figure 4: Prevalence of discovered codes in the stock market interviews



Note: This figure presents the frequency of codes discovered during our qualitative analysis of the stock market interviews (see Table A.2 for an overview of studies). Codes derived from hand-coding. For each code, we present the share of participants to which this code has been assigned at least once. For categories, we create a dummy equal to 1 if at least one code from the category was assigned to the interview, and zero otherwise.

Adding the task of researching companies, keeping up with earnings reports, and tracking market trends would be another layer of complexity and stress.

Many respondents perceive high recurring participation costs from owning stocks, due to the need for (i) research and (ii) monitoring. First, they believe that research is an essential task of an investor: “I know you are supposed to do research on stocks that you would like to invest in, for example if there was a recent oil spill, it would not be good time to buy stock in oil companies.” This research has to integrate a broad array of public information, such as historical price data,

financial accounts, earnings announcements, and analyst forecasts. As a result, respondents understand that “it takes time to research and understand a company’s product and how it fits into the marketplace” and anticipate “spending a couple of hours a day or maybe several hours 3–4 days a week online reading about market trends, emerging technologies or medicines, and companies on the uprise as well as those in decline.”

Yet, the perceived job requirements of an active investor continue beyond the point when the initial research is concluded and promising stocks have been identified, because one would “have to constantly monitor the current market condition and make sure that the stock [you] invested in would not suddenly crash.” Continuous monitoring is important because stock prices change “all the time, every day” and “if you sell [stocks] at a bad time you won’t get as much for the stock than if you sold them at a good time... and waiting for the peak time to sell seems best.” Many respondents thus seem to believe that making a return on one’s investments requires careful timing of the stock market, which only adds to the perceived costs of participation.

Collectively, these perceptions characterize what we refer to as the *active investing* model of stock market participation. This model is grounded in three fundamental beliefs. First, respondents believe that market inefficiencies allow investors to predict future stock prices through research and analysis. Second, they perceive stock prices as highly volatile, making continuous monitoring and quick decision-making important to avoid incurring large losses before prices adjust. Third, there is a shared view that investors make returns from short-term trades. These beliefs contribute to a view that participating in the stock market is a demanding, high-maintenance activity, which ultimately deters many respondents from investing.¹⁹

4.1.2 Mechanism 2: Knowledge barriers

Another prevalent theme among respondents is a concern that they lack fundamental knowledge about the stock market. Many respondents expressed uncertainty about the basic logistics of investing, such as how to get started, how to buy and sell stocks, or create an investment account. As one respondent stated: “I don’t know where to view the stock market, how people make money from stocks, and/or how they work. I don’t know how to invest, or what to start with when investing.” This perceived lack of knowledge extends beyond logistical aspects and into the complexity of the stock market itself:

To be honest, all of it sounds complicated. I barely know what any of the terms in the market mean. Like, I’m not familiar with dividends, mutual funds, etc. Aside from not having knowledge of the terms, I also don’t know what makes something a good investment vs a bad one. I wouldn’t even know where to start.

Many respondents point to the confusing jargon and technical concepts used in finance. This

¹⁹Duraj et al. (2025) independently discovers a similar active investing model after conducting 25 human-led qualitative interviews in Germany, providing further robustness to the finding from a different context. Furthermore, since their working paper was made public a year after our findings (Chopra and Haaland, 2023), it could not be the case that our AI interviewer “discovered” this hypothesis because it was part of the training data.

perceived need to grasp many complicated concepts before even beginning to invest leads many respondents to conclude that acquiring the necessary knowledge would be a time-consuming endeavor with uncertain rewards. As one respondent puts it, “it seems like there is a lot to learn, and that is a little intimidating.” Moreover, respondents voiced a deeper concern in the form of a belief that successful investing requires specialized knowledge and skills that they do not possess. Specifically, many respondents explain that successful investing requires an ability to predict stock prices, distinguish “good” from “bad” stocks, and effectively time stock trades ahead of everyone else. As one respondent explains:

I feel like I don't know which stocks are 'good' to purchase, or when a good time to purchase would be. Because of this, I feel like I would be flying blind and making decisions entirely at random, as compared to a more seasoned stock market veteran.

This perceived need for experience and expert knowledge both underscores the cost of stock market participation and reinforces the conclusion that without such knowledge, one is likely to incur losses. Another respondent articulated a similar concern, saying that they are “not quite sure how to read stocks and what patterns to look for,” and that “for this reason, I feel like if I were to try and invest in stocks without this knowledge I might lose money.” This combination of the perceived logistical complexity of investing, a belief in the necessity of extensive knowledge and experience, and the fear of losing money precisely because they lack those skills contribute to a perception of substantial entry barriers to investing.

4.1.3 Mechanism 3: Risk perceptions

A common theme that emerges when discussing the stock market are strong concerns about risks. For instance, many respondents express a deep-seated fear of losing a substantial part, or even all, of their “hard-earned” money on the stock market. Unlike risk-free assets where returns are guaranteed, respondents worry that they could end up with less money than they started with:

If I would lose any amount of money, no matter how large or small, it would make me sick to think about. I work so hard for everything I have, that knowing it could all be gone with the stock market failing would make me incredibly anxious.

A general unwillingness to make nominal losses, no matter how small, is frequently mentioned among participants, who are sometimes explicit that they would need a *guarantee* against making any losses to consider investing. Furthermore, the concern that an investment could “instantly lose all its value” reflects the belief that such catastrophic losses are not only possible, but likely, and could occur at any moment. The fear of catastrophic events leads many to conclude that the risks of investing are too large, and that owning stocks would therefore expose them to constant feelings of stress and anxiety because “the market changes everyday.”

The metaphor of gambling frequently surfaces in respondents’ accounts of the risks and returns of investing. This comparison is not merely metaphorical but reflects the mental model that the stock market operates with the same uncertainty and lack of control as a game of chance. For these respondents, the stock market is “like a lottery, where you can either win big or lose

everything” and it is “all a roulette of which stock will fail and which will win.” Investing in stocks is seen as being as risky and unpredictable as gambling, where the odds are perceived to be stacked against the average investor, supporting the view that the stock market is something for “when you have discretionary income to burn. . . not something you can rely on to build wealth.”

It is important to note that not all respondents see the stock market purely as a gamble. There is also a belief among respondents that the risks can be mitigated through knowledge and experience. As one respondent noted, investing in the stock market is risky “unless you know what you are doing.” Another respondent, who expresses a common concern about going “completely bankrupt,” puts it like this:

One of my biggest fears is going completely bankrupt since my stocks may fail. I also fear I may not know anything about the company I invest in and something huge happens where they close. I want to feel better about my experience and have that knowledge and funds.

These ideas tie back to previous themes, where respondents expressed that acquiring the necessary knowledge and skills could transform the stock market from a high-stakes gamble into a more manageable risk, implying that a seasoned investor faces a different risk-return trade-off. Yet, this knowledge comes at a considerable cost both in terms of time spent and, more importantly, the financial losses incurred during the costly learning process, which many feel they could not afford. As one respondent summarizes: “You could put money to buy stocks that end up losing value rather than gaining value and end up in a situation where you lost both time and money for participating.” This explains why many respondents prefer risk-free assets where neither time nor expertise are required to receive a guaranteed return on their investments.

4.2 Out-of-sample validation of the generated hypothesis

A key finding from our qualitative analysis of the stock market interviews is that many participants subscribe to an *active investing* mental model that makes stock market participation seem like a very substantial commitment of time and effort. If this mental model indeed acts as a barrier to stock market participation, it should be more prevalent among non-investors than among stock owners. To test this hypothesis, we conduct a large-scale survey in which we measure the prevalence of this mental model with tailored closed-end survey questions. A key benefit of this exercise is that it allows an out-of-sample replication of our qualitative findings which also addresses concerns about “overfitting” whereby spurious patterns in the interview corpus lead to chance discoveries that do not generalize.

Design For the out-of-sample validation survey, we recruited a balanced sample of 1,000 US investors and non-investors from Prolific in May 2024. Table A.3 presents summary statistics and Appendix Section F.1.6 contains the key instructions. We elicit stock market participation and provide a definition of the term. As a direct measure of the *active investing* mental model, we use agreement with the statement below, measured on a 5-point Likert scale:

“If you want to make any money in the stock market, you have to constantly monitor market conditions, keep up-to-date with all the news, and do a lot of research on the stocks that you plan to invest in.” (*Strongly disagree to strongly agree*)

In addition, we elicit other hypothesized facets of the *active investing* mental model. Specifically, we elicit participants’ beliefs about how often it would be “necessary to monitor market conditions and check one’s portfolio” on a 7-point scale from “less than once a year” to “multiple times a day,” their belief about the time per week required to manage one’s portfolio, whether it is important to react quickly to daily changes in market conditions and whether it is important to “actively pick a few stocks.” We also elicit the types of *public* information that respondents believe one would have to “regularly consult [...] to make a return on one’s portfolio,” such as earnings reports, company press releases, or social media. We deliberately chose very strong formulations for the questions, emphasizing that the described behaviors are of essential nature for successful stock market participation. We therefore interpret our measures as a conservative estimate of the true prevalence of the *active investing* mental model.


Results Column 1 of Panel A of Table A.6 shows that 76.9% of non-owners agree with the *active investing* mental model compared to only 55.1% of respondents who participate in the stock market, a difference of 39.6%. Non-investors are also 68.8% more likely to believe that one has to monitor one’s portfolio at least once a day (Column 2), 53.2% more likely to believe that quick reactions to daily changes in market conditions are necessary (Column 3), and 41.8% more likely to believe that selecting a few stocks that are likely to perform well is the key to successful stock market participation (Column 4). Moreover, non-investors select 36.8% of a standard deviation more public information sources as being required reading when owning stocks (Column 5). As a result, non-investors expect to spend 41.1% more time per week managing their portfolio compared to respondents who own stocks (Column 6). Panel B of Table A.6 shows that these results are not driven by differences in financial literacy, and are robust to controlling for a large set of background characteristics, including education, income, and financial assets. Panel C further shows that our main survey measure of the active investing model is strongly correlated with its subfacets, supporting the construct validity of the measure.

Taken together, these findings confirm that investors and non-investors have systematically different mental models of the stock market regarding the time and effort cost of investing. This confirms our key qualitative hypothesis from the analysis of our AI-led interviews, providing a proof of concept that AI-conducted interviews can generate hypotheses about economic behavior that can be empirically tested out-of-sample with quantitative methods.

4.3 Main qualitative findings from the tariff interviews

We now turn to our second application on economic policy views. An advantage of this setting is that we can compare reasoning between opponents and supporters of tariffs to see whether the

Figure 5: Frequency of discovered codes among supporters and opponents of import tariffs



Note: This figure presents the share of respondents that can be assigned a specific code based on the open-ended text responses they provided (in %). We use data from the tariff interview study (see Table A.2 for an overview of data collections). Codes assigned by AI coding. The vertical lines indicate the 10% code frequency thresholds. Support for US import tariffs elicited during the screener survey prior to the interview.

interviews uncover important differences in underlying reasoning.

In contrast to accounts questioning the ability of the American public to reason effectively about public policies (Bartels, 2005), our AI-led interviews reveal that both supporters and opponents of US import tariffs hold rather nuanced views about their effects (see Figure 5 for code frequencies).²⁰ Rather than being ignorant about the other side’s arguments, most respondents acknowledge the short-term negative effects of tariffs—such as rising prices—but differ in how they believe the economy will respond in the long term. What emerges are not simplistic opinions, but competing mental models that reflect divergent beliefs about how firms, consumers, and trade partners will respond to import tariffs.

²⁰A comprehensive discussion of the findings from the tariff interviews is beyond the scope of this paper. Instead of striving for narrative completeness in our discussion, we focus on more novel themes and factors that will be important for our validation and benchmarking exercises.

4.3.1 Opponents

When asked about their views on tariffs, opponents strongly focus on the many perceived short-term costs and long-term macroeconomic risks. They understand tariffs as being a “tax on the American consumer,” anticipating that firms will pass on tariff-induced costs to consumers:

The reason or reasons that I oppose the new 10% tariffs on imports is that the price increases are passed on to us, the consumer. In this current market, where prices keep increasing, people don't have extra money, and adding tariffs is going to make purchasing everything more expensive, making it more difficult for people to survive.

For many opponents, these concerns about inflation are compounded by worries about the regressive, distributional effects of tariffs. As one participant notes, tariffs “will affect every class of people, but it will hit the middle class, lower class the hardest.” On the supply side, opponents worry about rising unemployment as firms try to cut costs and, as one participant points out, tariffs will “likely cause many small and medium-size businesses to go out of business because of reduced sales, reduced profits, and uncertainty.” As a result, opponents unanimously expect tariffs to cause severe long-term economic damages, describing a possible “domino effect” where higher prices reduce consumer demand, force businesses closures, cause product shortages, and ultimately hurt growth. There is also considerable concern about deteriorating international relations, with opponents expecting other countries to “retaliate, plain and simple. They will not take the threat [of tariffs] lightly and they should not. It will isolate the US.” Taken together, the prevailing view among opponents is that tariffs are a self-inflicted wound motivated by political self-interest rather than public welfare.

When asked about potential benefits of tariffs, the most common response is that tariffs “only have negative effects.” Yet, even the most skeptical opponents understand what tariff supporters see as the goals of the policy: protecting jobs and reviving manufacturing. The disagreement lies not in the goals themselves, but in whether they are perceived to be achievable. Opponents are quick to point to many structural barriers that, in their view, make job creation unlikely. These include the time required to rebuild capacity (“this is a long, long process”), the absence of industrial infrastructure (“the factories aren’t built”), and broader labor market barriers. Specifically, factory work is seen as being unattractive and underpaid—“no one wants to work in factories and there is no way those jobs would pay a living wage”—and the US workforce is viewed as lacking both the training and the willingness to return to manual manufacturing jobs. Even if production came back, respondents expect limited employment gains due to automation. Ultimately, what many opponents invoke are competitive advantage arguments to explain that reshoring jobs would simply not make any sense economically:

The workforce in the US is largely, I mean, it's disproportionately white collar to blue collar compared to other countries. So that's why we use our competitive advantage to only manufacture high ticket items or to sell services rather than products. And other countries like China or Vietnam might have a lot of factories where people are willing to go in and work manual labor jobs, or, you know, jobs that create a product rather than just a service. So I see the need for us to take

advantage of, you know, our strengths and other countries to take advantage of their own strengths and for us to mutually benefit off of each other.

4.3.2 Supporters

Tariff supporters, for their part, begin from the same shared understanding of short-term costs—that prices will go up, that there may be short-term disruptions—but are much more optimistic about their potential long-term impacts. While the majority of opponents think tariffs have only negative effects, the large majority of tariff supporters have a “short-term pain, long-term gain” mental model of tariffs, where they expect long-term improvements in economic growth, fairness, and national strength.

Many participants describe tariffs as a means of restoring domestic self-sufficiency and reducing reliance on foreign imports. “We should be able to manufacture medical supplies, weapons, steel, etc. here in the US,” one participant argued, linking tariffs to both economic and national security. Others describe tariffs as tools to “push us to source our own goods from within the US and not other countries” or as part of a broader goal of “making everything in the U.S. and isolating the U.S. from the other parts of the world.” In this sense, tariffs are not only protective but a catalyst for restoring American jobs and independence by incentivizing the purchase of goods made in the US.

Closely linked to these policy goals are recurring themes of fairness and reciprocity in global trade. Many supporters feel that the US has long been treated unfairly, pointing to other countries’ high tariffs or exploitative labor practices as sources of competitive disadvantage. “For many years other countries charged tariffs on items imported from America, but we did not charge any tariffs to other countries,” one participant notes, concluding that “tariffs will assist our country in establishing a level playing field.” Another put it more bluntly: “If other countries are charging our country tariffs, then it only makes sense we do the same.” These fairness concerns often carry a moral undertone: imposing tariffs is not just economically sensible but “the right thing to do,” a way to ensure that “no country is better off than the other by trade.”

While the objectives of job creation and fairness in international trade relations are commonly mentioned, what the interviews uniquely reveal are the economic *mechanisms* that supporters invoke to explain why tariffs will work as intended.

First, many supporters believe that tariffs will *increase* innovation. They argue that removing access to cheaper imports will create incentives for US firms to innovate: “Necessity is the mother of invention,” one participant wrote: “Because tariffs will reduce the amount of foreign goods flooding US markets, American manufacturers will need to find ways to scale up their production and cut costs while preserving quality.” Others highlighted the need to develop more efficient production methods: “It will force companies to produce things locally. . . to build more advanced factories. . . which would involve a lot of innovation and change.”

Second, participants frequently expressed the view that consumer behavior will mitigate

the negative effects of price increases. “Prices of American-made goods would be lower than imported goods,” one respondent explained, “and knowing that their purchase would build our American economy would bring satisfaction as well.” Several described switching to domestic alternatives, avoiding non-essentials, or seeking higher-quality American products as natural consumer responses. This substitution channel is seen as buffering the cost of tariffs, reducing the real impact on household budgets through endogenous spending responses.

Finally, many supporters see tariffs not as permanent policies but as tactical instruments—a way to exert pressure on other countries and extract better trade deals. Several described tariffs as a “bargaining chip,” a “game of chicken,” or a means to “push [other countries] into better deals.” As one participant put it: “Initially, I think it would cause a lot of chaos and price spikes, but ultimately other countries would come to the negotiating table because they simply can’t maintain a 30% tariff.” While both supporters and opponents mention the risk of trade wars, their interpretation differs: for supporters, economic leverage and market size will allow the US to prevail in any retaliatory cycle.

4.3.3 Key findings

Taken together, these interviews show that disagreement over tariffs is less about ignorance and partisan cheerleading and more about disagreement about economic mechanisms. Supporters and opponents share the same views on short-term economic costs, but differ in how they believe economic actors will respond to shape the long-term effects of the tariffs. Opponents focus on structural barriers—missing infrastructure, labor shortages, or global competition—that they see as limiting the gains from tariffs. Supporters, by contrast, have a “short-term costs, long-term benefits” view in which they believe that firms will innovate, consumers will adapt, and the US can use its leverage to gain better trade terms. These competing mental models shape not only what respondents expect to happen, but also how they judge whether the policy is worth it.

5 Internal and predictive validity of AI-led interviews

We next address a natural question many economists used to studying revealed preference data might have: How meaningful are the responses obtained from AI-led interviews? Can we trust that the conversations give meaningful insights to predict and understand economic preferences and behaviors? Similar concerns were once raised about survey methods, with skepticism about whether self-reports could yield valid insights into preferences or behavior. We anticipate that this debate will be revisited in the context of AI-led interviews, where a potential concern is that interviewees just try to fill the conversation space created by the AI interviewer. While sociologists typically see the primary value of interviews as in generating novel ideas, economists frequently seek an additional output: measurements that can be used for quantitative analysis. This motivates the current section. We show that codes derived from AI-led interviews correlate with conceptually related survey measures (internal validity), and that they predict incentivized

behavior in a long-term follow-up (predictive validity).

5.1 Internal validity: Correlations with survey measurements


If participants provide meaningful and informative responses, we would expect a correlation between codes assigned to interview transcripts based on our coding manual and direct measurements of conceptually related constructs based on standard survey questions.²¹ To test this, both the stock market and the tariff interview studies included a battery of survey questions after the interview module. In the stock market study, we presented participants with a multiple choice list of common reasons for stock market non-participation in randomized order and asked participants to select all that apply. In the tariff study, we elicited expectations about the effects of imposing import tariffs using 5-point Likert scale questions (e.g., macroeconomic or distributive effects), which we standardize. For each study, we match survey questions with a theoretically matching code from our coding manual. For example, in the tariff case, we match the survey question eliciting beliefs about whether low-income households will gain or lose with the *regressive impact* code from our coding manual. We can match 14 (12) survey question from the stock market (tariff) studies with a code from our coding manuals. While we tried to include a comprehensive survey battery of structured Likert questions, the incomplete matching of coding illustrates that it is difficult to anticipate the relevant factors that will surface in interviews.

Figure 6 presents the results from separately regressing the survey measures on a dummy for the presence of their matched code. Correlations are sizable and in the expected direction. Panel A shows that participants are, on average, 45 p.p. more likely to select a reason for stock market non-participation from the multiple choice question if they have been assigned the matched code. Panel B shows that participants that are assigned a code score, on average, 0.52 standard deviations higher on the matched survey question. Figure 6 also presents estimates separately for subsamples defined by age, gender, and education. While we do not expect the prevalence of particular codes to be identical across demographic groups, we find no evidence that the strength of the relationship between interview-derived codes and their matched survey measures varies with participant characteristics. This suggests that AI-led interviews perform equally well in capturing economically relevant themes across different sociodemographic groups.

For the stock market interviews, we can additionally test whether interview codes are correlated with validated survey measures of economic traits, such as preferences, beliefs, and constraints. Table A.8 presents the results from this exercise. For example, Column 1 shows that participants who are assigned the *dislike of risk* code are 0.55 standard deviations more risk-averse using a validated risk preference measure (Falk et al., 2018). We find similarly pronounced associations for other economic factors, such as financial literacy (Column 2), subjective beliefs

²¹Our goal is to show that the measurements are correlated. We take no stance on which measurement is methodologically preferred or a more valid “ground truth.” For example, if researchers know precisely what they would like to measure, a tailored survey question might be better suited. If measurement should avoid priming a particular mechanism or if discovering unexpected mechanisms is key, an interview might be preferred.

Figure 6: Internal validity: Interview codes correlate with direct survey measurements



Note: This figure combines bivariate regression estimates from the stock market (Panel A) and tariff (Panel B) interview studies. We match post-interview multiple choice response options (Panel A) or 5-point Likert scale questions (Panel B) with a conceptually related code from our coding manual. In Panel A, we regress an indicator for having selected a specific multiple choice option (0/1) on a dummy for the presence of the code (0/1). In Panel B, we regress the standardized Likert scale response on a dummy for the presence of the code (0/1). In addition, each panel represents a pooled estimate when reshaping the data to the interview-by-survey measure level (“pooled estimates”), both in the full sample and subsamples defined by sociodemographics (gender, age, education). Labels in Panel A refer to the label of the multiple-choice option, while labels in Panel B refer to the matched code. Codes assigned by AI. 95% CIs derived from robust standard errors are shown.


(Column 3), measures of constraints (Column 4–5), or asset class preferences (Column 6). These findings demonstrate the internal and convergent validity of AI-led interviews.

5.2 Predictive validity: Behavior in a long-term follow-up

A complementary and demanding question specific to interviews with the goal of understanding economic behavior is whether participants’ implicitly or explicitly *stated* explanations during an AI-led interview are behaviorally predictive. To address this, we conduct a long-term follow-up with our stock market interview sample after eight months, including incentivized outcomes.

Design We invite stock market interview participants from Aug/Sep 2023 to a long-term follow-up survey in May 2024, and manage to recontact 70% of the original sample. Table A.3 provides summary statistics and Appendix Section F.1.4 contains the experimental instructions. A key design objective is generating testable behavioral predictions for a subset of codes from our coding manual. To focus on non-participation motives that only show up later during interviews, we selected four codes that are almost never assigned to initial responses: *active investing model*

Figure 7: Interview responses predict incentivized behavior eight months later



Note: Data from the long-term stock market interview follow-up (see Table A.2 for a study overview). Each panel separately shows the mean and standard error of an incentivized decision for participants that were either assigned a certain interview code from our coding manual, or not. Panel A: Belief about the share of actively managed funds that outperform passively managed funds. Panel B: WTP for a 3-month online investment course. Panel C: Share of respondents who only switch from a certain payment of \$1,000 in 12 months to investing \$1,000 in the S&P 500 for 12 months if offered full nominal insurance. Panel D: Share of respondents who are unwilling to invest in the S&P 500 even if offered full nominal insurance. p -values for t -tests for equality of means are shown.

(ACT, 0.5% frequency in the first response, see Figure 4), *high learning cost* (TML, 1.8%), *unwilling to make nominal losses* (UNL, 0.5%), *would never participate* (WNP, 0%).²² The survey includes three incentivized outcomes tailored to these interview codes. First, we elicit respondents' incentivized belief about the share of actively managed funds that yielded higher after-fee returns than passively managed funds in 2023 on a scale from 0 to 100%. Respondents could earn a bonus of \$2 if they guess within 2 p.p. of the actual value. This allows us to test whether respondents who revealed an “active investing” mental model during interviews believe that active portfolio management outperforms passive investing. Second, we elicited respondents' probabilistically incentivized willingness to pay (WTP) for three-month access to an “four-part series aimed at helping beginners understand the basics of stock market investments” from MasterClass, a platform providing high-quality learning content. Third, we elicited respondents' demand for loss protection in the stock market using a multiple price list. Respondents repeatedly choose between two options. Option A is a certain payment of \$1,000 in twelve months. Option B grants the value of \$1,000 invested in the S&P 500 in twelve months. Additionally, Option B caps the maximum loss from this investment (“loss protection”) at $X\%$, with X ranging from 0% to 100%. We implement a randomly selected choice for one out of every 200 respondents. This outcome provides two predictions. Respondents who reveal an “unwillingness to make nominal losses” (e.g., “I would need a *guarantee* on my investment”) should *only* switch away from the risk-free Option A if offered *full* insurance. In addition, people who emphasize that they would under “no circumstances” invest in stocks should choose Option A even if $X = 100\%$.

²²When regressing a dummy for participating in the follow-up survey on a set of dummies for whether these codes are assigned, we find no evidence of differential attrition. The p -value of a joint F -test is 0.416.

Results Panel A of Figure 7 shows that respondents with an *active investing* mental model report a 6.4 p.p. higher belief about the share of active funds that outperform the market compared to those who are not assigned the code ($p = 0.023$). Panel B shows that the average WTP for an investment course is 48% higher than the WTP among respondents who do not mention high learning costs as a barrier to investing in stocks ($p = 0.016$). Panel C shows that the probability of switching from the certain payment to an investment in the stock market *only if* offered full insurance against nominal losses increases by a factor of 2.8 if respondents reveal an unwillingness to invest in assets that include the possibility of “ending up with less than one started with” ($p = 0.002$). Finally, respondents assigned the *would never invest* code are more likely to prefer the certain payment over an investment with losses capped at zero ($p = 0.012$).

In sum, our results show that AI-led interviews capture internally valid and behaviorally predictive factors. This methodological result mitigates concerns about “cheap talk” or ex-post rationalization, and highlights the value of AI-led interviews for studying economic behavior.

6 Benchmarks for AI-led interviews

Having established that AI-led interviews can lead to novel insights across relevant economic settings, we now leverage the large scale of our data to answer methodological questions about how AI-led interviews compare to other data collection methods for gathering *qualitative* data. In Section 6.1, we introduce three progressively more demanding benchmarks, ranging from a single open-ended question to an “interview” without probing questions. In Section 6.2, we start with research applications aimed at measurement. We show that the property that human-led qualitative interviews measure “more” than open-ended survey questions carry over to AI-led interviews. In Section 6.3, we provide suggestive evidence that AI-led interviews are comparable to human-led interviews in their ability to efficiently extract information with probing question. Specifically, we show that AI-led interviews saturate in small samples typical for qualitative research studies.

6.1 Design of benchmarks

Survey We conducted the stock market benchmark survey ($N = 384$, Nov 2024) and the tariff benchmark survey ($N = 249$, July 2025) with US adults recruited from Prolific using the same screening criteria that we used in the corresponding interview studies (see Section 2.2 for the criteria and Table A.2 for study details). Both benchmark surveys use instructions identical to the interview surveys, except that we replace the AI-led interview with a total of six predetermined open-ended questions, presented on separate pages. To encourage effort, we frame these questions as the main part of the study when recruiting participants and ask for detailed responses (e.g., Section F.1.3). The first open-ended question in the benchmark surveys mirrors the opening question of the respective interview, such as “What are the main reasons you support or oppose the new 10 percent tariff on imports?” for the tariff benchmark. The

additional open-ended questions in the benchmark surveys are designed as “mirrors” to the topics explicitly mentioned in the topic guides of the AI interviews. For example, the second stock market interview topic asked the AI to “delve into the perceived barriers or challenges preventing them from participating in the stock market.” Our mirror question in the stock market benchmark survey thus becomes: “What barriers or challenges prevent you from participating in the stock market?” Section F.2.3 and F.1.3 present the wording of the mirror questions.

Benchmark 1: Single open-ended question Our first benchmark represents the *status quo* method for collecting open-ended data from study participants at scale (Haaland et al., 2025): We compare the codes discovered in response to the first open-ended question of the benchmark surveys with the codes discovered when using the full AI-led interview. This benchmark addresses the natural question whether researchers would overlook relevant factors that could affect their conclusions when restricting themselves to a stand-alone open-ended question.

Benchmark 2: Multiple open-ended questions Our second benchmark consists of the codes discovered when using the responses from all open-ended questions in the benchmark surveys. The mirror questions allow us to account for the role of the topic guide and the effect of simply asking more open-ended questions. A key remaining difference to an interview is that the benchmark surveys do not include dynamically generated probing questions.

Benchmark 3: Interviews without probing questions Our most demanding benchmark uses the transcripts of our AI-led interviews, but excludes the responses to the dynamically generated probing questions from the transcripts before assigning codes to participants. Compared to Benchmark 2, we additionally partial out the potential effects of (i) interacting with a conversation partner and (ii) using a chat interface.²³

6.2 Objective: Measurement and quantification in large samples


If the research objective is to measure and quantify the prevalence of different themes in *large samples*, we need to compare AI-led interviews to open-ended survey questions, as this is the only other method that can deliver qualitative data at scale. For this exercise, we define a code as *discovered* by a method if it would be assigned to at least 10% of participants.²⁴ We then compare the (i) share of discovered codes (extensive margin) and (ii) the distribution of code frequencies (intensive margin) across data collection methods.

Aggregate results Figure 8 presents the results. We start with the extensive margin. A single open-ended question would discover at most 21% of the codes discovered with the AI-led

²³This will also pick up positive effects of probing questions on respondents’ effort and engagement that arise from the feeling of having a genuine interaction—which would not arise without probing questions—thereby overestimating the effectiveness of this benchmark. We interpret this benchmark as a conservative upper bound.

²⁴We use this threshold for two reasons. First, AI coding is 6 p.p. more likely on average to assign a given code than human annotators (Section 2.3.2). Second, codes present in less than 4% of the data are unlikely to be noticed or included in the coding manual of a researcher focused on quantification.

Figure 8: AI-led interviews discover more major themes than open-ended survey questions




Note: This figure presents summary statistics for discovered and assigned codes across data collection methods and applications. We combine data from the stock market (Panel A) and tariff interviews (Panel B), and the corresponding benchmark surveys with pre-determined open-ended questions (see Table A.2 for an overview of studies). “Discovery rate” is the share of discovered codes relative to an AI-led interview based on a 10% frequency threshold. “Average number of codes” is the average number of unique codes assigned to a participant. For comparability, statistics are based on AI coding. We exclude infrequent codes that are assigned to less than 10% of participants in an AI-led qualitative interview. We report results for AI-led interviews and the three benchmarks defined in Section 6.1.

interview. Asking more open-ended questions raises this share to 35–54% depending on the application, implying that half of the theoretically discoverable codes remain unobserved. Even with our third benchmark—AI interviews without probing questions, which accounts for the interactive nature of the interview and the differences in response modes—40% to 50% of codes remain unobserved. This suggests that probing questions account for most of the gains of AI-led interviews over open-ended survey questions.

We observe similar patterns when turning to the number of codes assigned at the individual level. A single open-ended question yields at most 2.1 codes per participant on average. Asking more open-ended questions—either statically or as part of an AI-led interview without probing questions—increases the number of codes per participant. However, the number of themes per participant are about 50% lower compared to the 9 to 13 codes per participant in AI-led interviews. This suggests that one would strongly underestimate the richness of how people reason about the interview topic with open-ended survey questions without adaptive probing.²⁵

²⁵On average, Table A.7 shows that code frequencies are about 17 p.p. higher in AI-led interviews compared to

Figure 9: Stock market non-participation: AI-led interviews compared to survey benchmarks



Note: This figure presents the share of respondents that can be assigned a specific code based on the open-ended text responses they provided (in %). We combine data from the interview sample of our stock market interview study and the stock market benchmark survey (see Table A.2 for an overview of data collections). All participants answered open-ended questions about their reasons for not participating in the stock market. We report results for AI-led interviews and the three benchmarks defined in Section 6.1. All codes used for this exercise are derived from AI coding for comparability. The vertical lines indicate the 5% and 10% code frequency thresholds.

Taken together, this suggests that AI-led interviews reveal substantially more unique codes, themes, and factors than responses to standard open-ended survey questions. We next turn to the level of an individual code to characterize the *nature* of codes that remain hidden—unless we conduct an AI-led interview.

Code-level results The previous results mask considerable heterogeneity at the level of individual codes. We therefore turn to disaggregated results, focusing on the stock market interviews in Figure 9. Across all benchmarks, most codes remain below the 10% frequency threshold, with quantitative differences in code frequencies across benchmarks below 10 p.p. for most codes.

How would this affect the conclusions of a researcher interested in studying stock market non-participation? With a single open-ended question, most responses revolve around knowledge barriers (*LOK*) and income constraints (*NFI*). When asking multiple open-ended survey questions, complexity perceptions (*PCO*) and general risk aversion (*DOR*) would also surface as major themes. Against these benchmarks, the first broad conclusion is that researchers using an AI-led interviews would additionally discover conceptually distinct *types* of codes. For example,

multiple open-ended survey questions. Notably, the gains from additional open-ended questions skew towards the few “top-of-mind” codes stand-alone open-ended questions tend to elicit. If we exclude these codes, the average code frequency is low at 6.6% to 14%, even when asking multiple open-ended survey questions.

AI-led interviews would reveal codes tied to participants’ mental models (*ACT*, *SMG*, *ITM*) that rarely surface with open-ended survey questions, and participation cost codes related to the “active investing model” (*TML*, *TRM*). Similarly, codes related to social influences, such as peer effects (*PES*), the role of personal experiences (*PEX*), or emotions (*SAN*) only show up regularly in AI-led interviews. Importantly, AI-led interviews meaningfully change the order of code frequencies. For example, the third most common code in AI-led interviews—“fear of large losses” (60%)—is only assigned to 5% of participants based on a single open-ended survey questions. Similar patterns hold for social influences (*PEX*) or codes tied to emotions (*SAN*).

We replicate the above finding in the context of our import tariff interviews: AI-led interviews can reveal people’s underlying *mental models* and economic *mechanisms* that are largely absent from open-ended survey questions without adaptive probing (Figure B.5; see Appendix Section C.1 for details).


In sum, these findings suggest that interviews are not mechanically increasing code frequency, but can reveal highly prevalent factors that are simply not on top of people’s minds when responding to open-ended survey questions and that merely asking “more” open-ended survey questions is not an effective substitute for an AI-led interview. Consequently, the quantitative conclusions drawn from a data collection would likely to be different depending on whether we use AI-led interviews or open-ended survey questions.

Dynamics of discovery The above raises another natural question: When are codes discovered *within* an interview? In Appendix Section C.2, we analyze how the discovery of different themes unfolds over the course of an interview. We find that top-of-mind explanations, such as those focusing on economic constraints (e.g., low income), are revealed within the first few questions, while factors such as beliefs, mental models, and broader economic reasoning appear more gradually over time.

6.3 Objective: Saturation and hypothesis generation in small samples

We next focus on benchmarks for qualitative research applications of AI-led interviews. We have previously shown that AI-led interviews can generate novel hypotheses in Section 4, but how “efficient” are AI-led interviews at uncovering new themes compared to other alternatives? We study this question through the lens of *saturation*: How many interviews are required until additional data no longer yield substantively new themes? Intuitively, quicker saturation should correspond to higher information gain, suggesting that maximally informative probing questions are being asked. Practically, qualitative data collections are typically stopped once the last few interviews have not yielded new themes, providing an empirical stopping criterion for saturation. Saturation thus provides a metric to compare AI-led interviews with human-led interviews. We focus on inductive thematic saturation (Saunders et al., 2018; Guest et al., 2006), which human-led interview studies achieve after about 20 to 30 interviews. For example, in concurrent

Figure 10: AI-led interviews saturate in small sample sizes typical for qualitative research studies



Note: This figures presents data on inductive thematic saturation. We plot the expected share of codes discovered at progressively larger sample sizes (the curves) and the expected sample size at which a researcher would stop the data collection based on a saturation criterion (the markers). A code is defined as discovered at sample size N if at least two examples of the code have been observed. The stopping criterion for saturation is that no additional code was discovered based on the last five participants. Shares are calculated based on the set of codes the AI-led interview would have discovered in large samples, focusing on codes with at least 5% frequency. Expectations are calculated from $B = 500$ bootstrap repetitions that reshuffle the order in which data is observed. Panel A presents data from the stock market and Panel B from the tariff interview studies. The range for saturation in human-led interviews is represented by the shaded area at 20-30 interviews (Guest et al., 2006). We report results for AI-led interviews and the three benchmarks defined in Section 6.1.

work, Duraj et al. (2025) report saturation after 25 interviews on stock market non-participation and arrive at similar conclusions about the role of the *active investing* mental model, while Bustos et al. (2022) report saturation after 20 interviews with tax consultants.

We conduct a simulation exercise to estimate how fast AI-led interviews achieve saturation using a bootstrap with $B = 500$ repetitions that each randomly shuffle the order in which our original data was collected. For each sample size k , we define a code as *discovered* if the code is present in at least two interviews (to distinguish it from idiosyncratic statements; Saunders et al., 2018). We define the *stopping point* of the data collection as the smallest k such that the last five interviews did not discover any new codes. The share of discovered codes at sample size k is defined relative to the codes one could have discovered in the full sample. This exercise leverages the scale of our interviews, as researchers usually do not know the counterfactual of what would have happened if more interviews would have been conducted.²⁶

Figure 10 presents the results. The expected stopping point of AI-led interviews is a sample size of about $N = 25$ in both applications, at which point a researcher would on average have discovered 93% of all codes one could have observed in a sample of over $N = 380$. For

²⁶We assume that researchers can identify novel themes with perfect accuracy, in line with how saturation is conceptualized. This implies that the notion of “discovery” here is different from Section 6.2, where we require that a theme occurs at least 10% of interviews. In practice, this might be necessary not only to notice a code but also to prevent that too many themes are “discovered” based on anecdotes.

comparison, Guest et al. (2006) identified 94% of all possible codes with $N = 25$ human-led interviews. By contrast, our other three benchmarks would on average stop data collections too early based on a saturation criterion, thereby failing to discover both (i) individually important codes and a (ii) large share of possible themes. To illustrate (i), we focus on the *active investing* code from the stock market application: Researchers using AI-led interviews will discover the *active investing* code with a 94% probability when using the empirical saturation criterion to stop the data collection. Not asking probing questions lowers this chance of discovery to 75%, while only using multiple open-ended survey questions further decreases this to 52%. A single open-ended question would discover the code in only 4% of data collections.

In sum, the evidence suggests that AI-led interviews ask highly informative probing questions to saturate at a speed and at a level comparable to studies using human-led interviews that arrive at similar qualitative conclusions. This further strengthens the value of AI-led interviews for qualitative research purposes.

7 Conclusion

This paper introduces AI-led interviews as a novel method for collecting qualitative data at scale, allowing researchers to conduct unlimited interviews in parallel at costs similar to a survey. We run large-scale interviews in different economic contexts to answer four methodological questions: Are AI-led interviews of high quality? Do they generate novel insights into economic questions? Are interview responses internally and externally valid? And, finally, do AI-led interviews outperform other methods for collecting qualitative data at scale?

We show that the answers to these questions are positive. First, we document high interview quality on more than ten metrics grounded in best practices for qualitative interviewing, and find high participant engagement and satisfaction with the AI-led interview. Second, we show empirically that data from AI-led interviews allow hypothesis generation about important economic phenomena. Third, we find that interview responses have high internal validity and predict future behavior more than half a year later, which shows that AI-led interviews capture genuine and authentic concerns. Fourth, we show that AI-led interviews generate thematically richer responses than other survey methods, allowing researchers to more exhaustively capture participants' views both at the individual level and in the aggregate.

How can AI-led interviews be used productively in economics? First, AI-led interviews can be used to better measure and quantify how people reason about economic problems using their own words at scale, making factors visible that would be unobserved with open-ended survey questions. Importantly, the length of the interview is a choice variable: Researchers can opt for a "short interview" with a standardized open-ended question followed by two probing questions to reduce measurement noise by clarifying potential ambiguities, or flexibly increase the interview length if comprehensiveness is important, such as when studying high-dimensional

objects like mental models or political narratives. Second, AI-led interviews can be used to conduct interviews with all participants after a randomized controlled trial at low cost, making interview transcripts a source of data that can be used to quantitatively diagnose mechanisms behind treatment effects. Third, AI-led interviews can be used as a vehicle to deliver treatment protocols at scale, such as motivational interviewing (Chopra et al., 2026). Finally, AI-led interviews reduce the entry barriers to qualitative and mixed-methods research in economics. As AI-led interviews saturate quickly in small samples, it has never been easier to complement empirical work with qualitative evidence at all stages of the research process, ranging from hypothesis generation to designing, testing, and implementing a study.

References

- Andre, Peter, Ingar Haaland, Christopher Roth, Mirko Wiederholt, and Johannes Wohlfart, “Narratives about the Macroeconomy,” *Review of Economic Studies*, 2026, p. rdag014.
- , Philipp Schirmer, and Johannes Wohlfart, “Mental Models of the Stock Market,” *CESifo Working Paper*, 2024.
- Armantier, Olivier, Giorgio Topa, Wilbert van der Klaauw, and Basit Zafar, “An Overview of the Survey of Consumer Expectations,” *Economic Policy Journal*, 2017.
- Ash, Elliott, Ruben Durante, Marilia Grebenshchikova, and Carlo Schwarz, “DP16624 Visual Representation and Stereotypes in News Media,” Discussion Paper 16624, CEPR, Paris & London 2022.
- , Soumitra Shukla, and Jason Sockin, “Interviews,” *Working paper*, 2026.
- , Stephen Hansen, and Yabra Muvdi, “Large Language Models in Economics,” Discussion Paper DP19479, Centre for Economic Policy Research 2024.
- Baader, Malte, Kai Baron, Anna Hochleitner, and Jonas Pilgaard Kaiser, “The Stories We Tell: Societal Narratives and the Rise of the Far Right,” *mimeo*, 2025.
- Bartels, Larry M., “Homer Gets a Tax Cut: Inequality and Public Policy in the American Mind,” *Perspectives on Politics*, 2005, 3 (1).
- Bergman, Peter, Raj Chetty, Stefanie DeLuca, Nathaniel Hendren, Lawrence F. Katz, and Christopher Palmer, “Creating Moves to Opportunity: Experimental Evidence on Barriers to Neighborhood Choice,” *American Economic Review*, 2024, 114 (5), 1281–1337.
- Bewley, Truman F., “A Depressed Labor Market as Explained by Participants,” *American Economic Review*, 1995, 85 (2), 250–254.
- , *Why Wages Don’t Fall During a Recession*, Harvard University Press, 1999.
- Blinder, Alan, Elie R.D. Canetti, David E. Lebow, and Jeremy B. Rudd, *Asking about Prices: A New Approach to Understanding Price Stickiness*, Russell Sage Foundation, 1998.
- Boissin, Esther, Thomas H Costello, Daniel Spinoza-Martín, David G. Rand, and Gordon Pennycook, “Dialogues with large language models reduce conspiracy beliefs even when the AI is perceived as human,” *PNAS nexus*, 2025, 4 (11), pgaf325.
- Bordalo, Pedro, John Conlon, Nicola Gennaioli, Spencer Kwon, and Andrei Shleifer, “How people use statistics,” *Review of Economic Studies*, 2026, 93 (1), 250–285.
- Braghieri, Luca, Sarah Eichmeyer, Ro’ee Levy, Markus M. Mobius, Jacob Steinhardt, and Ruiqi Zhong, “Article-Level Slant and Polarization of News Consumption on Social Media,” *SSRN Working Paper*, 2024.
- Brynjolfsson, Erik, Danielle Li, and Lindsey Raymond, “Generative AI at Work,” *The Quarterly Journal of Economics*, 2025, 140 (2), 889–942.
- Bursztyrn, Leonardo, Georgy Egorov, Ingar Haaland, Aakaash Rao, and Christopher Roth, “Justifying dissent,” *The Quarterly Journal of Economics*, 2023, 138 (3), 1403–1451.
- Bustos, Sebastián, Dina Pomeranz, Juan Carlos Suárez Serrato, José Vila-Belda, and Gabriel Zucman, “The Race Between Tax Enforcement and Tax Planning: Evidence From a Natural Experiment in Chile,” Working Paper 30114, National Bureau of Economic Research 6 2022.
- Carlson, Jacob and Melissa Dell, “A Unifying Framework for Robust and Efficient Inference with Unstructured Data,” *arXiv preprint arXiv:2505.00282*, 2025.
- Charness, Gary, Brian Jabarian, and John A. List, “The next generation of experimental research with LLMs,”

Nature Human Behaviour, 2025, 9 (5), 833–835.

Chopra, Felix and Ingar Haaland, “Conducting Qualitative Interviews with AI,” *SSRN Working Paper No. 4572954*, 2023.

—, **Christopher Roth, and Johannes Wohlfart**, “Home Price Expectations and Spending: Evidence from a Field Experiment,” *American Economic Review*, July 2025, 115 (7), 2267–2305.

—, **Nicolas Roeber, Ingar Haaland, and Christopher Roth**, “Evaluating Behavioral Interventions at Scale with AI,” Discussion Paper 385, ECONtribute January 2026.

Corbin, Juliet and Anselm Strauss, *Basics of Qualitative Research: Techniques and Procedures for Developing Grounded Theory*, 4th ed., SAGE Publications, Inc, 2014.

Costello, Thomas H., Gordon Pennycook, and David G. Rand, “Durably reducing conspiracy beliefs through dialogues with AI,” *Science*, 2024, 385 (6714), eadq1814.

Creswell, John W., *Research designs. Qualitative, quantitative, and mixed methods approaches* 2009.

Cuevas, Alejandro, Jennifer V. Scurrell, Eva M. Brown, Jason Entermann, and Madeleine I. G. Daepf, “Collecting Qualitative Data at Scale with Large Language Models: A Case Study,” *Proceedings of the ACM on Human-Computer Interaction*, 2025, 9 (2).

de Quidt, Jonathan, Johannes Haushofer, and Christopher Roth, “Measuring and Bounding Experimenter Demand,” *American Economic Review*, 2018, 108 (11), 3266–3302.

Dell, Melissa, “Deep Learning for Economists,” *Journal of Economic Literature*, 2025, 63 (1), 5–58.

DeLuca, Stefanie, Lawrence F. Katz, and Sarah C. Oppenheimer, ““When Someone Cares About You, It’s Priceless”: Reducing Administrative Burdens and Boosting Housing Search Confidence to Increase Opportunity Moves for Voucher Holders,” *RSF: The Russell Sage Foundation Journal of the Social Sciences*, September 2023, 9 (5), 179–211.

Draca, Mirko and Carlo Schwarz, “How Polarised are Citizens? Measuring Ideology from the Ground up,” *The Economic Journal*, July 2024, 134 (661).

Duraj, Kamila, Daniela Grunow, Michael Haliassos, Christine Laudénbach, and Stephan Siegel, “Rethinking the stock market participation puzzle: A qualitative approach,” Technical Report, CESifo Working Paper 2025.

Durán, Rafael Jiménez, Karsten Müller, and Carlo Schwarz, “The Online and Offline Effects of Content Moderation: Evidence from Germany’s NetzDG,” *Working Paper*, February 2025.

Falk, Armin, Anke Becker, Thomas Dohmen, Benjamin Enke, David Huffman, and Uwe Sunde, “Global Evidence on Economic Preferences,” *Quarterly Journal of Economics*, 2018, 133 (4), 1645–1692.

Fuster, Andreas and Basit Zafar, “Survey Experiments on Economic Expectations,” in “Handbook of Economic Expectations,” Elsevier, 2023, chapter 4, pp. 107–130.

Galasso, Vincenzo, Tommaso Nannicini, and Debora Nozza, “We Need to Talk: Audio Surveys and Information Extraction,” *Working Paper*, 2024.

Geiecke, Friedrich and Xavier Jaravel, “Conversations at Scale: Robust AI-led Interviews with a Simple Open-Source Platform,” *SSRN Working Paper*, 2025.

Gentzkow, Matthew, Bryan Kelly, and Matt Taddy, “Text as data,” *Journal of Economic Literature*, 2019, 57 (3), 535–574.

Gibson, Kerry, “Bridging the digital divide: Reflections on using WhatsApp instant messenger interviews in youth research,” *Qualitative Research in Psychology*, 2022, 19 (3), 611–631.

Glaser, Barney and Anselm Strauss, *Discovery of Grounded Theory: Strategies for Qualitative Research*, Routledge, 2017.

Graeber, Thomas, Christopher Roth, and Constantin Schesch, “Explanations,” *CESifo Working Paper*, 2024.

—, —, and **Florian Zimmermann**, “Stories, Statistics, and Memory,” *The Quarterly Journal of Economics*, November 2024, 139 (4).

Grootendorst, Maarten, “BERTopic: Neural topic modeling with a class-based TF-IDF procedure,” *arXiv preprint arXiv:2203.05794*, 2022.

Guest, Greg, Arwen Bunce, and Laura Johnson, “How Many Interviews Are Enough? An Experiment with Data Saturation and Variability,” *Field Methods*, 2006, 18 (1), 59–82.

Guiso, Luigi and Paolo Sodini, “Household finance: An emerging field,” in “Handbook of the Economics of Finance,” Vol. 2, Elsevier, 2013, pp. 1397–1532.

Guyen, Semra Yuksel, Tobias Gårdhus, Andreas Bjerre-Nielsen, and Hjalmar Bang Carlsen, “Comparing AI-led to human-led chat-based interviews: motivations, initial results and challenges,” <https://osf.io/k7f89> 2025. Preprint. This manuscript has not been peer reviewed.

Haaland, Ingar, Christopher Roth, and Johannes Wohlfart, “Designing Information Provision Experiments,” *Journal of Economic Literature*, 3 2023, 61 (1), 3–40.

—, —, **Stefanie Stantcheva, and Johannes Wohlfart**, “Understanding economic behavior using open-ended

- survey data,” *Journal of Economic Literature*, 2025, 63 (4), 1244–1280.
- Hager, Anselm, Lukas Hensel, Christopher Roth, and Andreas Stegmann**, “Voice and political engagement: Evidence from a field experiment,” *Review of Economics and Statistics*, 2025, 107 (4), 1149–1158.
- Horton, John J.**, “Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus?,” Technical Report, NBER Working Paper 2023.
- Hübsch, Theresa, Robert Mahlstedt, Pia Pinger, Sonja Settele, and Helene Willadsen**, “Mental Models of High School Success,” Technical Report, University of Bonn and University of Cologne, Germany 2026.
- Hwang, Byoung-Hyoun, Don Noh, and Sean Seunghun Shin**, “How Investors Pick Stocks: Global Evidence from 1,540 AI-Driven Field Interviews,” SSRN Working Paper, HKU Jockey Club Enterprise Sustainability Global Research Institute December 2025.
- Jabarian, Brian and Luca Henkel**, “Voice AI in Firms: A Natural Field Experiment on Automated Job Interviews,” *SSRN Working Paper*, 2025.
- Jiménez-Durán, Rafael**, “The Economics of Content Moderation: Evidence from Hate Speech on Twitter,” *SSRN Working Paper*, 2021.
- Kamoi, Ryo, Sarkar Snigdha Sarathi Das, Renze Lou, Jihyun Janice Ahn, Yilun Zhao, Xiaoxin Lu, Nan Zhang, Yusen Zhang, Ranran Haoran Zhang, Sujeeth Reddy Vummanthala, Salika Dave, Shaobo Qin, Arman Cohan, Wenpeng Yin, and Rui Zhang**, “Evaluating LLMs at Detecting Errors in LLM Responses,” 2024.
- Khot, Tushar, Harsh Trivedi, Matthew Finlayson, Yao Fu, Kyle Richardson, Peter Clark, and Ashish Sabharwal**, “Decomposed Prompting: A Modular Approach for Solving Complex Tasks,” 2023.
- Knott, Eleanor, Aliya Hamid Rao, Kate Summers et al.**, “Interviews in the social sciences,” *Nature Reviews Methods Primers*, 2022, 2, 73.
- Kong, Aobo, Shiwan Zhao, Hao Chen, Qicheng Li, Yong Qin, Ruiqi Sun, Xin Zhou, Enzhi Wang, and Xiaohang Dong**, “Better Zero-Shot Reasoning with Role-Play Prompting,” 2024.
- Korinek, Anton**, “Generative AI for Economic Research: Use Cases and Implications for Economists,” *Journal of Economic Literature*, 2023, 61 (4), 1281–1317.
- Legard, Robin, Jill Keegan, and Kit Ward**, “In-depth Interviews,” in Jane Ritchie and Jane Lewis, eds., *Qualitative Research Practice: A Guide for Social Science Students and Researchers*, SAGE Publications, 2003, pp. 138–169.
- Link, Sebastian, Andreas Peichl, Christopher Roth, and Johannes Wohlfart**, “Attention to the Macroeconomy,” Available at SSRN 4697814, 2024.
- Ludwig, Jens and Sendhil Mullainathan**, “Machine Learning as a Tool for Hypothesis Generation,” *Quarterly Journal of Economics*, 2024, 139 (2), 751–827.
- Lusardi, Annamaria and Olivia S. Mitchell**, “Financial Literacy Around the World: An Overview,” *Journal of Pension Economics & Finance*, 2011, 10 (4), 497–508.
- Manski, Charles F.**, “Measuring Expectations,” *Econometrica*, 2004, 72 (5), 1329–1376.
- Mohajan, Haradhan Kumar**, “Qualitative research methodology in social sciences and related subjects,” *Journal of Economic Development, Environment and People*, 2018, 7 (1), 23–48.
- Morduch, Jonathan and Rachel Schneider**, *The financial diaries: How American families cope in a world of uncertainty*, Princeton University Press, 2017.
- Mutz, Diana C.**, *Winners and Losers: The Psychology of Foreign Trade*, Princeton University Press, 2021.
- Namey, Emily, Greg Guest, Amy O’Regan, Christine L. Godwin, Jamilah Taylor, and Andres Martinez**, “How Does Mode of Qualitative Data Collection Affect Data and Cost? Findings from a Quasi-Experimental Study,” *Field Methods*, 2020, 32 (1), 58–74.
- Nass, Clifford, Jonathan Steuer, and Ellen R. Tauber**, “Computers are social actors,” in “Proceedings of the SIGCHI Conference on Human factors in Computing Systems” 1994, pp. 72–78.
- Nelson, Laura K.**, “Computational Grounded Theory: A Methodological Framework,” *Sociological Methods & Research*, 2020, 49 (1), 3–42.
- Noy, Shakked and Whitney Zhang**, “Experimental evidence on the productivity effects of generative artificial intelligence,” *Science*, 2023, 381 (6654), 187–192.
- Peer, Eyal, David Rothschild, Ariel Gordon, Zak Evernden, and Ekaterina Damer**, “Data quality of platforms and panels for online behavioral research,” *Behavior Research Methods*, 2022, 54, 1643–1662.
- Puente, Manuela**, “Qualitative Insights on the Income-Health Gradient,” 2026.
- Pyatkin, Valentina, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi**, “Generalizing Verifiable Instruction Following,” 2025.
- Rubin, Herbert J. and Irene S. Rubin**, *Qualitative Interviewing: The Art of Hearing Data*, 3 ed., Sage Publications, 2011.

- Saunders, Benjamin, Julius Sim, Tom Kingstone, Shula Baker, Jackie Waterfield, Bernadette Bartlam, Heather Burroughs, and Clare Jinks**, “Saturation in qualitative research: Exploring its conceptualization and operationalization,” *Quality & Quantity*, 2018, 52 (4), 1893–1907.
- Shiller, Robert J.**, “Why do people dislike inflation?,” in “Reducing Inflation: Motivation and Strategy,” University of Chicago Press, 1997, pp. 13–70.
- Sivaramakrishnan, S., M. Srivastava, and A. Rastogi**, “Attitudinal factors, financial literacy, and stock market participation,” *International Journal of Bank Marketing*, 2017, 35 (5), 818–841.
- Small, Mario Luis and Jessica McCrory Calarco**, *Qualitative Literacy: A Guide to Evaluating Ethnographic and Interview Research*, University of California Press, 2022.
- Stantcheva, Stefanie**, “Understanding Tax Policy: How do People Reason?,” *The Quarterly Journal of Economics*, 2021, 136 (4), 2309–2369.
- , “How to Run Surveys: A Guide to Creating Your Own Identifying Variation and Revealing the Invisible,” *Annual Review of Economics*, 2023, 15, 205–234.
- Starr, Martha A.**, “Qualitative and mixed-methods research in economics: Surprising growth, promising future,” *Journal of Economic Surveys*, 2014, 28 (2), 238–264.
- Wuttke, Alexander, Matthias Aßenmacher, Christopher Klammer, Max M. Lang, Quirin Würschinger, and Frauke Kreuter**, “AI Conversational Interviewing: Transforming Surveys with LLMs as Adaptive Interviewers,” in Anna Kazantseva, Stan Szpakowicz, Stefania Degaetano-Ortlieb, Yuri Bizzoni, and Janis Pagel, eds., *Proceedings of the 9th Joint SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature (LaTeCH-CLfL 2025)*, Association for Computational Linguistics Albuquerque, New Mexico May 2025, pp. 179–204.
- Zhang, Yue, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen et al.**, “Siren’s Song in the AI Ocean: A Survey on Hallucination in Large Language Models,” *arXiv preprint arXiv:2309.01219*, 2023.
- Zhou, Jeffrey, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou**, “Instruction-Following Evaluation for Large Language Models,” 2023.

For online publication only:

Conducting Qualitative Interviews with AI

Felix Chopra

Ingar Haaland

March 9, 2026

Section A contains additional tables. Section B contains additional figures. Section C presents additional results. Section D contains additional details about AI-led interviews and a discussion of considerations for conducting AI-led interviews. Section E contains details about the coding of interview transcripts. Section F presents the key experimental instructions.

A Additional tables

Table A.1: Interview topic guide for the AI interviewer

	Individual economic behavior: Stock market interviews	Economic policy views: US import tariffs
Topic 1 (Opening question, pre-determined)	I am interested in learning more about why you currently do not own any stocks or stock mutual funds. Can you help me understand the main factors or reasons why you are not participating in the stock market? (6 questions)	First of all, what are the main reasons you support or oppose the new 10 percent tariff on imports? (6 questions)
Topic 2	Delve into the perceived barriers or challenges preventing them from participating in the stock market. (5 questions)	Counter-perspective prompt - invite them to consider the side that hasn't been emphasized yet (costs vs. benefits). Example opening question: "So far, we've mostly discussed [dominant theme] effects. Can you think of any potential [opposing theme] effects of tariffs?" (2 questions)
Topic 3	Explore a 'what if' scenario where the interviewee invests in the stock market. What would they do? What would it take to thrive? Probing questions should explore the hypothetical scenario. (3 questions)	Reasoning about tariffs - invite them to explain how they think tariffs work. Example opening question: "Imagine you were going to explain tariffs to someone unfamiliar with them and how they might affect different groups. How would you explain tariffs to them?" (4 questions)
Topic 4	Probe for conditions or changes needed for the interviewee to consider investing in the stock market. (2 questions)	What-if scenario - Suppose the United States suddenly raised tariffs to 30 percent on imports from all major trading partners. What do you think would happen, both initially and further down the road? (2 questions)
Topic 5	–	Wrap-up - drawing on the full conversation, clarify unresolved ambiguities and deepen promising themes that deserve more exploration. (3 questions)
Closing questions (pre-determined)	(1) As we conclude our discussion, are there any perspectives or information you feel we haven't addressed that you'd like to share? (2) Reflecting on our conversation, what would you identify as the main reason you're not participating in the stock market?	Is there anything you would want to know or see that could change your mind about the tariffs?

Note: This table provides an overview of the interview structure and topic guide for our applications (see Section 2.2 for details). The table provides the *exact instructions* from the AI interviewer's prompt on which aspects to focus at each stage of the interview. These instructions are provided to both the topic agent and the probing agent. We indicate the instructed number of questions per topic in bold. The first question of the interview ("opening question") and the closing questions of the interview are identical across participants.

Table A.2: Overview of data collections

<i>Study and population</i>	<i>Recruitment</i>	<i>n</i>	<i>Study description</i>
Stock market interviews			
General population (US), non-stockowners	Prolific, Aug-Sep 2023, screener survey for recruitment	381	Descriptive survey with an integrated 30-minute text-based interview led by an AI interviewer. Post-interview survey questions about the interview experiences. Additional survey modules on factors associated with stock market non-participation, respondent characteristics and household finances.
Long-term interview follow-up survey			
Stock market interview participants	Prolific, May 2024	266	Long-term follow-up with participants from the stock market interview study. Descriptive survey that includes incentivized outcomes to study whether selected interview codes predict choices: (i) demand for insurance against nominal losses, (ii) willingness to pay for an investment course, (iii) incentivized belief about the performance of actively managed funds.
Out-of-sample validation survey			
General population (US)	Prolific, May 2024	1,000	Descriptive survey with equal share of stockowners and non-owners. Survey questions aimed at measuring whether respondents hold the “active investing” mental model that we hypothesized based on the qualitative analysis of the interview transcripts from the stock market interviews.
Selection experiment			
General population (US)	Prolific, Nov 2023	499	Respondents are asked to select all studies that they would like to participate in if we run it: (i) 40-minute survey, (ii) 40-minute interview with a human interviewer, (iii) 40-minute interview with an AI interviewer. These studies are shown on separate pages in randomized order.
Stock market benchmark survey			
General population (US), non-stockowners, recruitment identical to stock market interviews	Prolific, Nov 2024	384	Respondents are asked to answer six pre-determined open-ended survey questions about their reasons for stock market non-participation. The first question is identical to the opening question of our AI-led interviews; the other open-ended questions mimic the topic guide of our interviews.
Tariff interviews			
General population (US)	Prolific, May-Jun 2025, screener survey for recruitment	385	Descriptive survey with an interview with an AI chatbot about US import tariffs from April 2025. Post-interview survey questions about the interview experience. Additional survey questions about tariffs.
Tariff benchmark survey			
General population (US), recruitment identical to tariff interviews	Prolific, July 2025	249	Survey with six pre-determined open-ended survey questions about their views on tariffs, following the topic guide of the tariff interview study.

Table A.3: Summary statistics for the stock market studies

	(1) Interview sample	(2) Long-term follow-up	(3) Benchmark survey	(4) Validation survey	(5) Selection study
Age	39.36 (14.03)	42.24 (13.87)	34.35 (11.48)	37.15 (12.59)	42.63 (14.97)
Female	0.64 (0.48)	0.62 (0.49)	0.69 (0.46)	0.50 (0.50)	0.49 (0.50)
College degree	0.53 (0.50)	0.56 (0.50)	0.52 (0.50)	0.58 (0.49)	0.72 (0.45)
Full-time employment	0.47 (0.50)	0.52 (0.50)	0.37 (0.48)	0.46 (0.50)	0.52 (0.50)
Caucasian/White	0.77 (0.42)	0.79 (0.41)	0.65 (0.48)		
African American/Black	0.12 (0.33)	0.13 (0.33)	0.13 (0.34)		
Hispanic	0.11 (0.31)	0.11 (0.31)	0.14 (0.35)		
Northeast	0.18 (0.38)	0.17 (0.37)	0.21 (0.41)	0.21 (0.41)	0.26 (0.44)
Midwest	0.24 (0.43)	0.21 (0.41)	0.21 (0.41)	0.18 (0.38)	0.22 (0.41)
West	0.17 (0.38)	0.18 (0.38)	0.21 (0.41)	0.25 (0.43)	0.15 (0.36)
South	0.41 (0.49)	0.44 (0.50)	0.37 (0.48)	0.36 (0.48)	0.38 (0.48)
Household income (thousands)	71.91 (38.76)	69.29 (36.91)	72.94 (38.42)	76.07 (53.94)	84.70 (53.01)
Financial assets (thousands)	77.98 (138.65)	82.71 (145.38)	55.34 (114.32)	104.81 (162.19)	
Non-mortgage debt (thousands)	27.67 (60.06)	22.73 (51.22)	29.75 (54.31)	22.18 (47.37)	
Homeowner	0.51 (0.50)	0.52 (0.50)	0.38 (0.49)	0.46 (0.50)	
Home value (thousands)	243.01 (138.07)	239.49 (135.65)	312.84 (148.93)		
Mortgage debt (thousands)	124.29 (91.28)	112.82 (78.64)	156.96 (128.89)		
Two months of liquid savings	0.60 (0.49)	0.61 (0.49)	0.55 (0.50)	0.68 (0.47)	
Observations	379	266	384	1,000	499

Note: This table presents the mean and standard deviation (in brackets) of a range of background variables for the stock market studies (see Table A.2 for an overview of studies), as indicated in the column title. “Home value” and “Mortgage debt” are conditional on non-zero values. Empty cells indicate that the variable was not collected as part of the study.

Table A.4: No differential selection into AI interviews compared to surveys

	Dependent variable: Demand for study participation (binary)					
	(1)	(2)	(3)	(4)	(5)	(6)
AI interview	0.005 (0.034)	-0.022 (0.060)	0.007 (0.035)	0.003 (0.049)	0.007 (0.049)	-0.010 (0.075)
AI interview x Willing to write		0.036 (0.072)				
AI interview x Stock non-participation			-0.014 (0.116)			
AI interview x Above median age				0.004 (0.068)		
AI interview x Female					-0.012 (0.067)	
AI interview x College						0.009 (0.084)
N	333	333	333	333	332	332
R ²	0.000	0.003	0.008	0.000	0.001	0.007
Mean of dep. var.	0.895	0.895	0.895	0.895	0.898	0.898

Note: This table shows regression estimates using data from the selection experiment (see Table A.2 for an overview of studies). The dependent variable is a dummy for whether respondents would like to participate in the described study. All regressions restrict to the first study described to respondents, using vignettes that described an AI-led interview or a survey study. “AI interview” is a dummy for whether the study was described as including an AI-led interview, or not. “Willing to write” is a dummy for writing at least 20 words in a screener survey. “Stock non-participation” is a dummy for non-investors. “Above median age” is a dummy for having above median age. “Female” is a dummy for female respondents. “College” is a dummy for having completed a college degree. Robust standard errors are shown in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A.5: Heterogeneity in interview experience across sociodemographics

	Dependent variable:				
	(1) Interview satisfaction (std.)	(2) Naturalness (std.)	(3) Preference for AI-led interviews (std.)	(4) Interest in another interview (0/1)	(5) Interview duration (minutes)
Age (above median)	0.178** (0.072)	0.323*** (0.071)	0.091 (0.073)	0.011 (0.013)	5.396*** (1.082)
Female	0.045 (0.075)	0.108 (0.074)	0.105 (0.074)	0.013 (0.015)	-3.452*** (1.109)
College degree	-0.091 (0.073)	-0.009 (0.073)	-0.109 (0.073)	0.009 (0.014)	-1.559 (1.079)
N	766	766	766	766	766
R ²	0.010	0.030	0.008	0.005	0.070
Mean of dep. var.	0.000	-0.000	0.000	0.965	29.197
Interview fixed effect	Yes	Yes	Yes	Yes	Yes

Note: This table shows regression estimates using pooled data from the stock market interview study and the tariff interview study (see Table A.2 for an overview of studies). The dependent variables are satisfaction with the interview (Column 1), the perceived naturalness of the conversation (Column 2), the stated preference for an AI-interviewer over a human interviewer (Column 3), interest in another interview (Column 4), and the interview duration (Column 5). We winsorized the interview duration at the 5th and 95th percentiles, and we standardize Likert scale responses to have a mean of zero and a standard deviation of one within each study. Each column presents estimates from regressing the dependent variables on a set of dummy variables. All regressions include study fixed effects. Robust standard errors in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A.6: Out-of-sample replication of the active investing mental model in a large-scale survey

	Dependent variable:					
	(1) Active trading mental model	(2) Daily monitoring important	(3) Quick reactions to market important	(4) Actively picking winner stocks important	(5) Public info required to read (z-scored)	(6) Log weekly portfolio management time
<i>Panel A.</i>						
Non-owner	0.218*** (0.029)	0.132*** (0.027)	0.173*** (0.031)	0.207*** (0.030)	0.368*** (0.062)	0.411*** (0.073)
Constant	0.551*** (0.022)	0.192*** (0.018)	0.325*** (0.021)	0.495*** (0.022)	-0.182*** (0.043)	0.575*** (0.053)
N	1,000	999	999	999	999	999
R ²	0.053	0.023	0.031	0.045	0.034	0.031
<i>Panel B.</i>						
Non-owner	0.176*** (0.032)	0.120*** (0.031)	0.106*** (0.033)	0.155*** (0.034)	0.334*** (0.069)	0.334*** (0.082)
Financial literacy (0-3)	-0.113*** (0.018)	-0.072*** (0.019)	-0.124*** (0.020)	-0.089*** (0.018)	-0.070* (0.040)	-0.225*** (0.043)
N	999	999	999	999	999	999
R ²	0.098	0.051	0.099	0.089	0.045	0.086
Controls	Yes	Yes	Yes	Yes	Yes	Yes
<i>Panel C.</i>						
Active trading model		0.307*** (0.024)	0.387*** (0.028)	0.301*** (0.033)	0.783*** (0.065)	1.210*** (0.068)
Financial literacy (0-3)		-0.040** (0.018)	-0.082*** (0.020)	-0.059*** (0.019)	0.010 (0.040)	-0.093** (0.040)
N		999	999	999	999	999
R ²		0.139	0.219	0.148	0.151	0.291
Controls		Yes	Yes	Yes	Yes	Yes

Note: Data from the out-of-sample validation survey for the stock market interviews, conducted in May 2024 (see Table A.2 for a study overview). The dependent variables in columns 1–4 are binary variables. “Non-owner” is a dummy for respondents without any exposure to the stock market. “Active investing mental model” is a dummy for agreement with the active investing view. “Daily monitoring important” is a dummy for respondents selecting “daily” or “multiple times a day” when asked about how often it is necessary to monitor market conditions and check one’s portfolio. “Quick reactions to market important” is a dummy for selecting “the key to stock market success is to react quickly to daily market changes.” “Actively picking winner stocks important” is a dummy for selecting “the key to stock market success is to actively pick a few stocks that will increase in value and avoid stocks that will decrease in value.” “Public info required to read” is a z-scored number of public information sources that respondents think one has to “regularly consult when participating in the stock market to make a return on one’s portfolio.” “Log weekly portfolio management time” is the log of the weekly hours that respondents think one would have to spend managing one’s portfolio, measured on an interval scale from “less than 1 hour per week” to “more than 15 hours per week.” All regressions in Panels B and C include the following set of sociodemographic controls: age, gender, log household income, and indicators for having completed a college degree, working full-time, owning a house, having a mortgage, and having non-zero savings. “Financial literacy” is the number of correctly answered questions from the “Big 3” financial literacy inventory. Robust standard errors are shown in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A.7: AI-led interviews discover more themes than open-ended survey questions

	AI-led interview	Survey benchmark		
		AI-led interview: No probing	Multiple open-ended questions	Single open-ended question
	(1)	(2)	(3)	(4)
A. Economic behavior (stock market)				
Discovered codes	29	15	10	6
Average number of codes per participant	9.1	4.4	3.9	2.1
Average frequency of codes	31.3%	15.1%	13.4%	7.3%
... excluding top-of-mind codes	25.0%	9.0%	6.6%	2.7%
Discovery rate	100.0%	51.7%	34.5%	20.7%
B. Economic policy views (tariffs)				
Discovered codes	37	22	20	6
Average number of codes per participant	13.1	7.1	6.7	2.0
Average frequency of codes	35.5%	19.1%	18.2%	5.4%
... excluding top-of-mind codes	32.3%	15.1%	14.0%	2.9%
Discovery rate	100.0%	59.5%	54.1%	16.2%

Note: This table presents summary statistics for discovered and assigned codes across data collection methods and applications. The table combines data from the stock market (Panel A) and tariff interviews (Panel B), and the corresponding benchmark surveys with pre-determined open-ended questions (see Table A.2 for an overview of studies). Columns correspond to different data collections or benchmarks, as defined in Section 6.1. For comparability, statistics are based on AI coding. We exclude codes assigned to fewer than 10% of participants in an AI-led interview. “Discovered codes” is the number of codes that can be assigned to at least 10% of participants. “Average frequency of codes” is derived by computing the share of participants that can be assigned a given code, and then averaging this number across codes. “Average frequency of non-top-of-mind codes” is analogously derived by excluding codes that can be assigned to more than 10% of participants when asking only a single open-ended question. “Discovery rate” is the share of discovered codes relative to Column 1.

Table A.8: Interview codes are correlated with direct measurements of economic preferences, beliefs, and constraints


	(1)	(2)	(3)	(4)	(5)	(6)
	Risk-taking (Falk et al. 2018) (std.)	Big 3 financial literacy (std.)	Subjective belief: % chance of financial loss	Log assets	Age (in years)	Homeowner (binary)
DOR: Dislike of risks	-0.547*** (0.101)					
LOK: Lack of knowledge		-0.259** (0.113)				
FLL: Fear of losses			5.901** (2.446)			
NFI: No funds				-0.924*** (0.236)		
TOR: Too old/retired					23.001*** (1.838)	
PTA: Prefer tangible assets						0.219** (0.102)
Constant	0.319*** (0.079)	0.188** (0.096)	34.686*** (1.637)	10.499*** (0.200)	37.919*** (0.697)	0.044*** (0.011)
N	381	381	381	354	381	381

Note: This table shows regression estimates using data from the stock market interview study (see Table A.2 for an overview of studies). The regressors are binary variables for whether a code (e.g., “Dislike of risks”) was assigned to the participant based on the hand-coding of interview transcripts, or not. The dependent variable in column 1 is the willingness to take risks (Falk et al., 2018), which we standardize such that larger values denote higher willingness to take risks. The dependent variable in column 2 is the standardized number of correctly answered financial literacy questions from the “big 3” financial literacy inventory (Lusardi and Mitchell, 2011). The dependent variable in column 3 is the subjective belief about the probability of negative returns on the US stock market over the next 12 months. Robust standard errors in parentheses.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.


B Additional figures

Figure B.1: Validation of AI coding: Code frequencies derived from AI vs human coding across 12,192 decisions and 32 codes




Note: This figure presents a scatterplot of code frequencies derived from human coding against those derived from AI coding, using the interview transcripts from the stock market interview study (see Table A.2 for a study overview). In particular, we construct a dummy variable at the participant level that takes value one if a given code was assigned at least once, and zero otherwise. Across all interview codes, the mean difference between AI coding and human coding is 6.4 p.p. The global agreement between the AI and human coders is 85.6% across 12,192 interview-code pairs.

Figure B.2: Tariff interviews: Additional participant perceptions of questions




Note: This figure uses survey data from the tariffs interview study (see Table A.2 for an overview of studies). Each panel plots the distribution of responses to different Likert scale questions about the AI-led interview. “Open-endedness of questions” asked respondents to evaluate the open-endedness of the questions. “Neutrality of questions” asked respondents to evaluate the neutrality of the questions. “Comfort during interview” is participants’ comfort in interacting with the AI interviewer during the interview.

Figure B.3: The timing of discoveries over the course of the stock market interviews




Note: This figure uses hand-coded interview transcripts from the stock market interview study. We plot the discovery process $d_c(t)$, i.e., the share of participants that can be assigned a specific code c based on all of their responses to the first t interview questions. Topic transitions are indicated as vertical lines. Each curve represents the discovery process for a separate code (see Figure 4). Selected codes are highlighted.

Figure B.4: AI-led interviews increase response length compared to open-ended survey questions



Note: This figure shows the mean number of characters written in response to a question in an AI-led interview and in response to an open-ended survey question of equivalent nature. For this comparison, we combine data from the interview studies and the benchmark surveys (see Table A.2 for an overview of data collections). From interviews, we keep the opening question and the first question from each interview topic. These questions were designed to have mirror questions in the benchmark survey. We exclude the predetermined closing questions. Panel B splits the sample by whether participants used the voice input method at least once (any audio) or decided to answer questions exclusively by typing on their keyboard (no audio). 39% of tariff interviewees use the voice input method at least once. See Section 6.1 for details. Standard error of the mean shown as error bars.

Figure B.5: Tariff interviews: AI-led interviews compared to other survey benchmarks



Note: This figure presents the share of respondents that can be assigned a specific code based on the open-ended text responses they provided (in %). We combine data from the tariff interview study and the corresponding benchmark survey (see Table A.2 for an overview of data collections). “AI-conducted interviews” presents code frequencies based on the full AI-led interview transcripts from our main experiment (i.e., whether a given code was assigned at least once based on the full interview). “Multiple open-ended questions” presents code frequencies based on the responses to six open-ended survey questions in our benchmark survey. These open-ended questions correspond to the topic guide of the main experiment and thus cover the same topics as our main interview, except for any follow-up questions that an interviewer might have asked. “Single open-ended question” presents code frequencies based *only* on the response to the first open-ended question in our benchmark survey. All codes used for this exercise are derived from AI coding for comparability. The vertical lines indicate the 5% and 10% code frequency thresholds.

C Additional results

C.1 Benchmarking tariff interviews: Code-level results

Starting with tariff opponents, Figure B.5 shows that responses to a single open-ended question would reveal the “higher consumer prices,” “tax on consumers” and “only negative effects” codes. While asking additional open-ended questions elicits some additional concerns, only the AI-led interviews reveals key mechanisms that explain *why* opponents are skeptical about potential benefits of tariffs (e.g., “barriers to employment effects,” “competitive advantage logic”). Among tariff supporters, we similarly find that key mechanism-based arguments are largely absent from survey responses to open-ended questions. For example, without an AI-led interview, researchers might struggle to detect reasoning related to innovation (“more innovation”) or consumer substitution (“behavioral changes mitigate neg. effects”) that explain *why* supporters are not as concerned about the short-term costs of tariffs. Similarly, using data from open-ended survey questions alone would suggest that supporters almost exclusively see positive, while opponents almost exclusively see negative aspects of tariffs. Yet, when conducting a conversation, both sides express awareness of the key arguments posited by the other side; in other words, the opposing views on tariffs revolve much more around *disagreement* about which mechanisms will dominate in the long term rather than ignorance about the other side’s views. Taken together, these results underscore that AI-led interviews not only elicit a broader set of themes but are essential for uncovering how individuals reason about economic policies, and why they come to different conclusions.

C.2 The timing of discoveries

Our findings from Section 6 that AI-led interviews can lead to the discovery of themes and recurring patterns that would be largely unobserved in text responses to a single open-ended survey question raise a natural follow-up questions about the timing of discovery: How long does it take to reveal a previously unobserved factor within an interview? What is the marginal return to asking *another* interview question? How does the discovery process differ across codes?

To address these questions, we leverage the fact that the interview transcripts are hand-coded at the response level. We derive the share $d_c(t)$ of respondents that can be assigned code c based on their cumulative responses to the first t interview questions, but not taking into account information that was revealed later in the interview. We refer to $d_c(t)$ as the *discovery process* for code c . Figure B.3 plots the discovery process of all codes that are assigned to at least 5% of respondents by the end of the interview. For the top-of-mind codes *Lack of knowledge* and *Low income*, the first six probing questions capture most of the discoveries, with marginal returns diminishing beyond that point. In contrast, *Dislike of risks*, another top-of-mind response, shows a more gradual discovery process that only plateaus near the final questions. These patterns

of gradual discovery are even more pronounced for the “high-gain” codes, such as *Gambling model* or *Active investing model*, where the marginal returns from further probing remain high throughout the interview. The discovery process for *Active investing model* even shows a steep rise after the third topic transition, where the interviewer introduces a “what if” scenario involving stock market investment. Similar jumps for other codes around the topic transitions demonstrate the value of a topic guide that explores the same underlying questions from multiple angles.

C.3 Text production

A common issue with open-ended survey responses is brevity that makes text classification difficult: a sizable share of responses cannot be classified due to ambiguity and lacking context, rendering the data unusable for downstream analysis. We hypothesize that AI-led interviews are more engaging surveys, thereby increasing response length to otherwise equivalent questions. To test this, we compare the average response length in the AI-led interviews with the response length in our benchmark surveys, focusing only on questions for which the benchmark surveys include a “mirror” question (i.e., we exclude probing questions).

Figure B.4 shows that AI-led interviews increase the response length on a per-question basis by 56% in the stock market interviews (Panel A) and by 39% in the tariff interviews among participants that did not use the option to answer questions by recording a voice message (Panel B).¹ However, among the 39% that responded by voice, “text” output increases by 173%. However, Appendix Section C.4 shows that longer audio messages are as rich as typed responses, consistent with prior research on the role of response modes (Namey et al., 2020). Finally, we find that text production is very stable over the course of the entire AI-led interview, as shown in Figure 2.

C.4 Interview mode: Voice vs typed input

During our import tariff interviews, participants could respond either by typing or by recording a voice message. While this choice is endogenous, it offers interesting correlational evidence on the relevance of the response mode. 39% used the voice message functionality at least once, and The time participants take completing the interview is virtually identical with 27 minutes for both input modes on average, and we detect no difference in interview attrition rates across input modes ($p = 0.165$). As expected, people can talk faster than they can write: text production is 112% higher with voice input compared to the 3,430 characters participants typed on average (p -value < 0.001). Perhaps surprisingly, but mirroring the findings from randomized trials (Namey et al., 2020), this substantial difference in text does not necessarily translate into richer responses: the number of unique codes assigned from our coding manual is only 8% higher ($p = 0.009$)

¹Note that participants could not copy-paste text from AI tools such as ChatGPT into the chat interface due to programmatic restrictions added for the tariff interviews, which rules out that these results are driven by AI use. For the stock market interviews, self-reported weekly usage of AI tools is uncorrelated with text production ($p = 0.815$).

with voice input (14.3 codes) compared to typing (13.1). Moreover, in unreported regressions, we find that the frequency of individual codes is largely similar across modes. A notable exception is that voice input users are more likely to mention specific examples, corroborating the findings from Galasso et al. (2024). While response mode will surely matter in some contexts and for specific participant population, it is reassuring that the practical differences might be muted in economic applications where participants can actively choose the mode they prefer.

D Additional details about AI-led interviews

D.1 Prompts

We conducted our interviews with the frontier model available at the time. In 2023, we use the gpt-4-0613 and gpt-4-0314 versions of the GPT-4 model family for the stock market interviews. In 2025, we use gpt-4o-2024-08-06. The stock market interviews made use of the probing, summary, and topic agents. The tariff interviews used the probing, topic and moderator agents of our modular framework (see Section 2 for a description of these agents).

We provide the complete prompt templates for our agents as part of our GitHub repository (see below). These templates include {placeholders} that are replaced programmatically to reflect the current state of the interview. For example, {summary} is replaced with the summary agent’s current synopsis of the conversation on previous interview topics if the “summarization” option is activated. Otherwise, it is replaced with the full and non-summarized conversation (see Section 2 for a description of this optional feature). Similarly, the placeholder {current_topic_history} is replaced with the verbatim questions and answers belonging to the currently investigated interview topic. These messages are inserted into the prompt using the convention “Interviewer: ... Interviewee: ...”.

For prompts, additional implementation details and model configurations, we refer to our open-source repository with setup instructions: <https://github.com/fchop/interviews>

E Additional details on the coding of interviews

E.1 Human coding

This section contains additional information about the hand-coding of interview transcripts. Four trained research assistants from the University of Copenhagen and the NHH Norwegian School with field experience as interviewers or experience working in qualitative research projects, were given the task of hand-coding the interviews according to various coding schemes.

Training Research assistants read a handbook chapter on interviewing methods (Legard et al., 2003) and were provided with detailed coding manuals that included the definition of each variable along with several examples of what type of questions or answers would fall into the

different categories. Research assistants received extensive training before starting with the task in the form of multiple rounds of double and triple coding, followed by group discussions about how to reconcile coding differences until a joint interpretative understanding of all codes was reached. The probability of agreement for a randomly selected code is 83% across research assistants.

Thematic coding We randomly assigned interviews to research assistants. For each code, research assistants indicated their level of confidence in the code (high vs low). Research assistants assigned codes to individual responses to interview questions. A code was assigned to the k^{th} response if the conversation up until that point included sufficient unambiguous evidence to warrant assigning the code. If in doubt, research assistants were instructed to wait until the code could be unambiguously assigned at a later point during the interview when more evidence had accumulated. Each response could be assigned multiple codes.

Stock market interviews We conducted two separate rounds of hand-coding for all stock market interview transcripts, which include about 12,000 questions and responses. The first round of hand-coding focused on the nature and properties of questions asked by the AI interviewer and the subsequent responses (e.g. whether questions are open-ended or responses display signs of discomfort; see Section 3.2). The second round of hand-coding, conducted after the first round of coding, was based on the coding manual derived from the inductive thematic analysis and focused on coding participants' reasons for not participating in equity markets. A direct benefit of conducting two rounds of coding is that research assistants had already immersed themselves in the interview transcripts before starting with the more demanding task of coding respondents' reasons for not participating in the stock market. When we discussed the corresponding coding manual for the first time, the research assistants already had a robust understanding of many codes and could immediately recognize the concepts from interviews they had read previously. This ensured a high shared conceptual understanding of corresponding codes. The research assistants were also told that creating a coding scheme is an iterative process and were encouraged to suggest potential changes to the coding scheme for stock market non-participation reasons based on their readings of the interviews.

Tariff interviews For tariff interviews, 200 out of 384 interview transcripts were randomly selected and coded for the analyses shown in Section 3. This amounts to about 6,000 questions and responses. This annotation was done after the annotation of the stock market interviews.

E.2 AI coding

We use the following prompt template for AI coding. The {placeholders} in the template are replaced with information from the coding manual, and the interview transcript is appended at the end of the prompt.

Context: You are an expert in coding qualitative interviews with a strong background in sociology and

qualitative research methods. You are given an interview transcript containing numbered question-answer pairs exploring why the respondent does not participate in the stock market.

Task: Your task is to decide whether there is clear evidence of the code 'code' using a systematic three-step approach. Follow the steps below carefully.

Step 1: Extract Evidence

- *Identify the ****strongest evidence**** from the interview transcript that directly relates to the definition of code.*
- *Reference specific answers from the transcript (e.g., question 3).*
- *Include supplementary evidence only if it provides essential context or clarification.*
- *If no relevant evidence exists in the transcript, explicitly state this.*
- *****Do not introduce any information not present in the transcript.*****

Step 2: Evaluate Evidence

- *Determine if the extracted evidence is ****unambiguous**** according to these criteria: ****Assign '1' if:**** - The evidence directly matches the definition of code without making further assumptions.*
- *The respondent explicitly states personal experiences or beliefs that align with code.*
- *****Assign '0' if:**** - The evidence is ambiguous or fits another code better.*
- *Conflicting evidence weakens the case for code.*
- *****Handling Overlaps with Other Codes:*****
- *If the evidence aligns better with another code, exclude it from code and justify the alternative classification.*

Step 3: Justify Your Decision

- *Provide a clear justification for your classification.*
- *Reference the extracted evidence explicitly.*
- *Explain why the evidence meets or does not meet the criteria for code.*
- *If assigning '0', clarify why the evidence is ambiguous or better fits another code.*

Coding Guidelines:

1. *****Evaluate All Responses:**** Review every question-answer pair for full context.*
2. *****Consider Related Codes:**** Before confirming, check if the evidence aligns more closely with another code.*
3. *****Combine Relevant Evidence:**** Include all answers that collectively support or refute the code.*
4. *****Conservative Approach:**** Assign '1' only if the evidence unambiguously meets the definition.*
5. *****Context Matters:**** Supplementary answers can provide context but cannot create evidence.*
6. *****Use Examples for Guidance Only:**** Refer to the examples to understand the definitions, but do not rely on them when extracting evidence. Base your analysis solely on the interview transcript without introducing information from the examples.*

Full Coding Scheme for Reference: {list of all codes}

Specific Guidance on code: {definition}

Examples (for reference only): {examples}

Response Format:

Return the following structured response and nothing else:

```
{ "primary_evidence": "[The specific answer from the below interview transcript that provides the strongest evidence.]",  
  "supplementary_evidence": "[Additional context from the interview transcript, or 'None'.]",  
  "classification": [1 or 0],  
  "justification": "[A brief explanation that justifies the classification based on the extracted evidence.]" }
```

Interview Transcript: {transcript}

F Experimental instructions

This section contains the core experimental instructions for our studies. Text in *italics* is not part of the instructions. Table A.2 provides an overview of the data collections.

F.1 Stock market studies

F.1.1 Screener survey (August 2023)

The next question is about the following problem. In questionnaires like ours, sometimes there are participants who do not carefully read the questions and just quickly click through the survey. This means that there are a lot of random answers which compromise the results of research studies. To show that you read our questions carefully, please choose both “Extremely interested” and “Not at all interested” as your answer in the below question.

Given the above, how interested are you in sports? [Extremely interested, Very interested, A little bit interested, Almost not interested, Not at all interested]

Only respondents who select (a) and (e) pass the attention check.

Please explain: What is your opinion about the yearly switch to daylight saving time? Do you like or dislike it? Please use about 15-30 words.

Only respondents who write at least 20 words are eligible to participate in an interview.

What was your family’s gross household income in 2022 in US dollars? [Dropdown menu]

1. How do you currently save or invest your money? Please check all that apply. [Response options: (a) I am currently unable to save money; (b) Savings or checking account; (c) Certificates of Deposit (CDs); (d) Bonds or bond funds; (e) Money market accounts; (f) Individual stocks or stock mutual funds; (g) Real estate or property; (h) Retirement accounts (e.g. 401k, IRA); Other]

2. Do you currently own any individual stocks or stock mutual funds? [Response options: Yes, I own individual stocks.; Yes, I own stock mutual funds.; Yes, I own both individual stocks and stock mutual funds.; No, I do not own any individual stocks or stock mutual funds.]

3. Do you plan to buy individual stocks or stock mutual funds over the next three months?

4. Do you have a financial advisor who manages your investments for you? [Yes, No]

5. Do you own any individual stocks or stock mutual funds indirectly, such as through pension savings, retirement accounts, or other investment vehicles? [Response options: Yes, I own stocks or stock mutual funds indirectly.; No, I do not own any stocks or stock mutual funds indirectly.]

Respondents who selected (a) and not (f) in Question 1; “No” in Question 2; “No” in Question 3 and 4; and “No” in Question 5 were invited the interview survey.

F.1.2 Stock market interview survey (August/September 2023)

How often do you typically use AI tools such as ChatGPT, Bard or Bing per week? [Several times a day, About once a day, A few times a week, Once a week, Less than once a week, Never]

(New survey page)

Please tell us about yourself.

What is your age? [Dropdown menu]

What is your gender? [Male, Female]

Which category best describes your highest level of education? [Dropdown menu]

What is your current employment status? [Dropdown menu]

(New survey page)

Next up: The Interview. You will now participate in an interview with an AI chatbot that will ask you questions, similar to how a human interviewer would. The topic of the interview is stock market participation. The AI chatbot is informed that you indicated to us in a previous survey that you do not

currently own any stocks or stock mutual funds. The interview will take about 20 minutes.

How it works: The interview uses a chat interface, much like texting on your phone. When a question appears, you can type your response. After your response, it takes between 5 to 10 seconds for the next question to appear. Take your time in crafting your responses - there is no rush. And don't worry, there are no right or wrong answers.

Please note: Your responses are entirely anonymous and will be used solely for research purposes. Please do not reveal any personally identifiable information such as your name. The AI chatbot is instructed not to give financial advice or explain financial concepts. Kindly avoid such queries.

Try it out: Below is an example of the chat interface. Please try it out before you proceed. *[Example interface shown with mock questions]*

(New survey page)

To ensure that the interview is as meaningful as possible, it's important that your responses provide sufficient detail. This means going beyond short statements, and diving into the "why" and "how". We are interested in your unique perspective and voice.

So, as we move forward, please aim to write at least two to three full sentences in your answers. This will provide enough depth and context for a rich and meaningful conversation. Of course, if you don't have an opinion on a topic or you're unsure, it's completely fine to express that too, with responses such as "I don't know".

(New survey page)

Let's start the interview. Please write detailed and meaningful responses to each question to the extent that you are comfortable sharing. There are no right or wrong answers as we are interested in your unique experiences and perspectives. It can take between 5 to 10 seconds until the next question appears. Feel free to use this pause to reflect on the conversation. *[Chat interface]*

(New survey page)

Here is a summary of the previous interview. *[Text box with a summary of the interview]*

Does the above summary accurately represent your views expressed in the interview? [Yes, No]

(Only if "No" was selected in the previous question:) What is inaccurate in the above summary?

(New survey page)

Thank you for completing the interview with the AI chatbot!

Before this study ends, we have some traditional survey questions relating to your financial behavior and beliefs as well as a few demographic questions that we want to ask you.

(New survey page)

What are your main reasons for not participating in the stock market? Please check all that apply.

- a) I don't know how to get started
- b) Wealth too small to invest
- c) Only few years until retirement
- d) Need cash on hand / liquidity
- e) Risk of losing a substantial part of investments
- f) Uncertainty of returns
- g) Don't trust financial professionals
- h) Stock market is too complex
- i) Setting up an investment account is too costly
- j) Managing a portfolio is too costly (time, effort, money)
- k) Don't like to think about my finances
- l) Lack of knowledge
- m) Negative experiences with the stock market
- n) Stock market returns are not high enough
- o) My labor income is too uncertain

- p) Family, friends, or coworkers advise against stocks
- q) Prefer more tangible assets
- r) Intended to invest in stocks but never got around to it
- s) None of the above

Except for “None of the above”, we randomize the order in which items are presented.

(New survey page)

How would you rate your overall experience with the interview conducted by the AI chatbot? [Extremely positive, Very positive, Somewhat positive, Neither positive nor negative, Somewhat negative, Very negative, Extremely negative]

How natural did the conversation with the AI chatbot feel? [Extremely natural, Very natural, Somewhat natural, Neutral, Somewhat unnatural, Very unnatural, Extremely unnatural]

If you were to participate in a future study involving a qualitative interview conducted through a chat interface, would you prefer texting with the same AI chatbot or an actual human interviewer? [Strongly prefer AI chatbot, Moderately prefer AI chatbot, Slightly prefer AI chatbot, No preference, Slightly prefer human interviewer, Moderately prefer human interviewer, Strongly prefer human interviewer]

Would you be interested in participating in an interview with an AI chatbot again? [Yes, No]

(New survey page)

In this question we present eight different scenarios for US stock market returns over the next 12 months.

Please let us know how likely you think it is that each scenario will occur. Please type in the number to indicate the probability, in percent, that you attach to each scenario. The probabilities of the eight scenarios have to sum up to 100%.

The US stock market return over the next 12 months will be...

- a) More than 30% : _____
- b) Between 20% and 30% : _____
- c) Between 10% and 20% : _____
- d) Between 0% and 10% : _____
- e) Between -10% and 0% : _____
- f) Between -20% and -10% : _____
- g) Between -30% and -20% : _____
- h) Less than -30% : _____

Assume that your household's car broke down and the repair costs \$1,000. How difficult would it be for your household to take out a loan to finance this repair? [Very easy, Easy, Neither easy nor difficult, Difficult, Very difficult]

In case of an unexpected decline in income or increase in expenses, does your household have at least two months of income available in cash, bank accounts, or easily accessible funds?

(New survey page)

What is the current value of all financial assets (cash, savings accounts, bonds, stocks, etc.) that your household owns? Please include holdings through retirement accounts. [Dropdown menu]

What is the current value of all non-mortgage debt (credit card, student debt, medical debt, auto loans, etc.) that your household owes? [Dropdown menu]

Do you own or rent your current main residence? [Own, Rent, Other]

Do you have any mortgage debt outstanding against your household's main residence? [Yes, No]

(New survey page)

(Only shown if the respondent has outstanding mortgage debt:) What is the current value of all mortgage debt (incl. home-equity based borrowing, HELOCs, etc.) that is outstanding against your household's main residence? [Dropdown menu]

(Only shown for homeowners:) What do you think is the current value that your household's main residence would sell for? [Dropdown menu]

Suppose you had \$100 in a savings account and the interest rate was 2% per year. After 5 years, how much do you think you would have in the account if you left the money to grow? [More than \$102, Exactly \$102, Less than \$102, Don't know]

Imagine that the interest rate on your savings account was 1% per year and inflation was 2% per year. After 1 year, with the money in this account, would you be able to buy... [More than today, Exactly the same as today, Less than today, Don't know]

Do you think the following statement is true or false? Buying a single company stock usually provides a safer return than a stock mutual fund. [True, False, Don't know]

Which of the following statements is correct? If somebody buys the stock of firm B in the stock market, then... [He owns a part of firm B., He has lent money to firm B., He is liable for firm B's debts., None of the above statements are correct., Do not know.]

In general, how willing or unwilling are you to take risks? [11-point scale from "0 (completely unwilling to take risks)" to "10 (completely willing to take risks)"]

How willing are you to give up something that is beneficial for you today in order to benefit more from that in the future? [11-point scale from "0 (completely unwilling to do so)" to "10 (completely willing to do so)"]

Generally speaking, would you say that most people can be trusted or that you need to be very careful in dealing with people? [You can't be too careful, Most people can be trusted]

(New survey page)

Which of the following best describes your race or ethnicity? [Dropdown menu]

Are you of Hispanic, Latino, or Spanish origin? [Yes, No]

How many people, including yourself, currently live in your household? [Dropdown menu]

How many children do you have? [Dropdown menu]

What is your region of residence? [Dropdown menu]

In which state do you currently reside? [Dropdown menu]

(New survey page)

Thank you for participating in our survey! If there are any remarks that you would like to make or clarifications that you would like to obtain, please do let us know by writing them into the field below. [Text entry field]

F.1.3 Stock market benchmark survey (November 2024)

We conducted a screener survey identical to the screener survey from August 2023. Eligibility to participate in the benchmark survey was then determined based on identical criteria as in the main survey from August 2023.

In a previous survey, you told us that you are currently not participating in the stock market. That is, you do not own any stocks or stock mutual funds. On the next pages, we will ask you six open-ended questions about your reasons for not participating in the stock market.

It's important that your responses provide sufficient detail. Please aim to write at least two to three sentences in your answers. If you don't have an opinion on a topic or you're unsure, you can express that with responses such as "I don't know".

There are no right or wrong answers; we are simply interested in your unique experiences and perspectives.

(New survey page. The following open-ended questions are shown on separate pages each.)

We are interested in learning more about why you currently do not own any stocks or stock mutual funds. Can you help us understand the main factors or reasons why you are not participating in the stock market?

What barriers or challenges prevent you from participating in the stock market?

Imagine you invest in the stock market. What actions would you take, and what do you think would help you succeed?

What conditions or changes would need to happen for you to consider investing in the stock market?

As we conclude, are there any perspectives or information you feel haven't been addressed that you'd like to share?

Reflecting on our conversation, what would you identify as the main reason you're not participating in the stock market?

F.1.4 Long-term interview follow-up survey (May 2024)

We include a standard attention check after the consent form and then elicit basic demographic characteristics: age, gender, education, employment, and income. We measure these as in our main survey. We do not replicate these instructions.

Demand for insurance against losses

(continued)

An opportunity to get exposure to the US stock market

We will now give you the opportunity to decide between investing \$1,000 in the US stock market for 12 months or getting \$1,000 for sure in 12 months.

Loss protection

Loss protection allows you to shield your stock market investments from losses. How much you are shielded depends on the size of the loss protection. **The higher the loss protection, the less money you have at risk.** For example:

- With 0% loss protection, all your money is at risk. You can lose \$1,000 of the initial \$1,000 investment.
- With 50% loss protection, only half your money is at risk. You can maximally lose \$500 of the initial \$1,000 investment.
- With 100% loss protection, none of your money is at risk. You cannot lose a single dollar of your initial \$1,000 investment.

If the stock market increases, you still get 100% of the upside no matter the level of your loss protection.

How much loss protection do you require to put your money at risk?

You will repeatedly choose between two options that involve payments to your Prolific account in **12 months**.

Both options are **worth \$1,000 today**. The two options differ in how their value might change over the next 12 months:

- **Option A:** Will have a guaranteed return of **0%** (and thus be worth \$1,000 in 12 months).
- **Option B:** Will provide the same return as the US stock market over the next 12 months (as measured by the S&P 500 index) with varying levels of loss protection.

How do payments work?

1 out of every 200 survey participants will be randomly selected. If you are among the selected participants, we will randomly select one of the choices you make below and implement it.

- If Option A is implemented, we will transfer a \$1,000 bonus to your Prolific account in 12 months.
- If Option B is implemented, we will transfer \$1,000 adjusted for the return of the US stock market over the next 12 months and the amount of loss protection. We will also send you a monthly update about the value of your investment.

	Which option do you prefer?	
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 0% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 10% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 20% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 30% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 40% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 50% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 60% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 70% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 80% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 90% loss protection.
\$1,000 with a guaranteed 0% return.	☐	\$1,000 invested in the stock market with 100% loss protection.



Willingness to pay for an investment course

(continued)

How much do you value a stock market investment course?

MasterClass is an online education platform known for its wide range of professionally made courses with top experts from various fields.

They've recently introduced a four-part series aimed at helping **beginners understand the basics of stock market investments** and **learn how to successfully navigate the market**. As one of the instructors puts it, "Wall Street makes things very complicated, but it doesn't have to be that way."

How are choices made?

You will repeatedly choose between two options:

- **Option A:** You get free 3-month access to the "Mastering the Markets" stock market investment course from MasterClass.
- **Option B:** Varying amount of bonus payments to your Prolific account.

1 out of every 25 survey participants will be randomly selected. If you are among the selected participants, we will randomly select one of the choices you make below and implement it.

► Click **here** if you would like to know how you would receive access.

Which option do you prefer?

Stock market investment course	☐	☐	\$0 bonus
Stock market investment course	☐	☐	\$2 bonus
Stock market investment course	☐	☐	\$4 bonus
Stock market investment course	☐	☐	\$6 bonus
Stock market investment course	☐	☐	\$8 bonus
Stock market investment course	☐	☐	\$10 bonus
Stock market investment course	☐	☐	\$12 bonus
Stock market investment course	☐	☐	\$14 bonus
Stock market investment course	☐	☐	\$16 bonus
Stock market investment course	☐	☐	\$18 bonus
Stock market investment course	☐	☐	\$20 bonus



Belief about the performance of actively managed funds

Your beliefs about the performance of actively managed funds

A stock market fund pools money from multiple investors to invest in a diversified portfolio of stocks.

It is common to make a distinction between active funds and passive funds:

- Active funds are managed by professional fund managers who actively trade stocks and aim to outperform the market by **picking superior stocks and monitoring market conditions**.
- Passive funds try to obtain the average market return without actively picking stocks or taking market conditions into account when buying and selling.

Every year a large investment research company measures how many percent of active funds were able to outperform their passive counterpart (after fees).

How many percent of U.S. active funds do you think outperformed their passive counterpart in 2023? If your answer is within ± 2 percentage points of the actual answer, you will receive a **\$2 bonus**.

I think % of active funds outperformed their passive counterpart.

Note: We will tell you the actual answer and whether you will receive the \$2 bonus at the end of the survey.



F.1.5 Selection experiment (November 2023)

We elicit similar sociodemographic variables as in the Screener Survey (August 2023). We do not replicate these instructions here.

Thank you for completing the background questions in this survey. Based on your responses, you might be eligible to participate in future potential studies.

Would you be interested in participating in future studies? On the next few pages, you will find a brief description of each study. For each study described, please indicate if you are interested in receiving an invitation to participate in that study.

AI interview (order randomized)



A 40-minute study with a qualitative interview

By Felix Chopra

£10.00 • £15.00/hr 40 mins 80 places Limited capacity

In this study, you will participate in a text-based interview with an AI and answer some standard survey questions about your beliefs and behavior. We expect the survey to take about 40 minutes in total.

Devices you can use to take this study:

☐ Desktop

If we run this study, would you like to participate in it?

☐ Yes

☐ No

>>

Human interview (order randomized)



A 40-minute study with a qualitative interview

By Felix Chopra

£10.00 • £15.00/hr 40 mins 80 places Limited capacity

In this study, you will participate in a text-based interview with a human and answer some standard survey questions about your beliefs and behavior. We expect the survey to take about 40 minutes in total.

Devices you can use to take this study:

☐ Desktop

If we run this study, would you like to participate in it?

☐ Yes

☐ No

>>

Survey version (order randomized)



A 40-minute survey

By Felix Chopra

£10.00 • £15.00/hr 40 mins 80 places Limited capacity

In this study, you will participate in a survey and answer some standard survey questions about your beliefs and behavior. We expect the survey to take about 40 minutes in total.

Devices you can use to take this study:

☐ Desktop

If we run this study, would you like to participate in it?

☐ Yes

☐ No

>>

F.1.6 Validation survey (May 2024)

We elicit similar sociodemographic variables as in the Screener Survey (August 2023).

To what extent do you agree or disagree with the following statement? “If you want to make any money in the stock market, you have to constantly monitor market conditions, keep up-to-date with all the news, and do a lot of research on the stocks that you plan to invest in.” [Strongly disagree, Disagree, Neither agree nor disagree, Agree, Strongly agree]

To succeed with stock market investments, how often do you think it is necessary to monitor market conditions and check one’s portfolio? [Less than once a year, Once a year, Once a month, Once a week, Multiple times a week, Once a day, Multiple times a day]

To succeed in the stock market, how much time do you think you would typically need to spend every week to manage your investments? [Less than 1 hour per week, 1–3 hours per week, 4–6 hours per week, 7–10 hours per week, 11–15 hours per week, More than 15 hours per week]

Which of the two statements below best describes your views?

- a) The key to stock market success is to react quickly to daily market changes
- b) The key to stock market success is to ignore daily market changes.

Which of the two statements below best describes your views?

- a) The key to stock market success is to actively pick a few stocks that will increase in value and avoid stocks that will decrease in value.
- b) The key to stock market success is to invest broadly in all stocks without making an active selection.

Which of the following types of information do you think one would have to regularly consult when participating in the stock market to make a return on one’s portfolio? Select all that apply.

- a) Earnings reports (e.g., quarterly or annual financial statements of companies)
- b) Market reports (e.g., stock market performance summaries and trends)
- c) Forecasts from stock market analysts (e.g., buy/sell recommendations and target prices for individual stocks)
- d) Company press releases (e.g., new product launches, management changes)
- e) Technical analysis (e.g., charts and patterns predicting stock price movements)
- f) Insider trading reports (e.g., transactions by company executives)
- g) Social media and forums (e.g., investment communities, Twitter, Reddit)
- h) None of the above

F.2 Tariff studies

F.2.1 Screener survey (May 2025)

Before starting, we have another quick check to ensure you’re a real person who carefully reads survey instructions. Bots and automated or inattentive responses compromise our research, so thank you for your help! To confirm you’re human and paying close attention, please ignore your actual state of mind for the next question.

Instead, follow the instructions below and select three response options, namely: Select the option that includes the word "relaxed." Select the third option in the list. Select the first option listed.

Which of the following describes your current state of mind? [I feel focused, I feel tired, I feel happy, I feel low, I feel relaxed, I feel stressed, None of the above, I prefer not to answer]

(New survey page)

What is your opinion about the yearly switch to daylight saving time? Do you like or dislike it? Please use about 15 to 30 words. [text entry field]

(New survey page)

Do you approve or disapprove of the Trump administration substantially increasing tariffs on goods imported from most countries that trade with the US? [Strongly approve, Somewhat approve, Somewhat disapprove, Strongly disapprove]

(New survey page)

Please think about the increased trade of goods and services between the US and other countries in recent decades. How do you think increased trade has affected the US? [The US gained more than it has lost from increased trade, The US gained about as much as it has lost from increased trade, The US lost more than it has gained from increased trade]

(New survey page)

On economic policy matters, where do you see yourself on the liberal/conservative spectrum? [Very liberal, Liberal, Moderate, Conservative, Very conservative]

(New survey page)

In politics, as of today, do you consider yourself a Republican, a Democrat, or an Independent? [Republican, Democrat, Independent]

Who did you vote for in the 2024 presidential election between Donald Trump and Kamala Harris? [Donald Trump, Kamala Harris, Someone else, Did not vote]

(New survey page)

(Displayed if respondent identifies as Republican) Do you consider yourself a strong Republican or a weak Republican? [Strong Republican, Weak Republican]

(Displayed if respondent identifies as Democrat) Do you consider yourself a strong Democrat or a weak Democrat? [Strong Democrat, Weak Democrat]

(Displayed if respondent identifies as Independent) Do you lean towards the Republican Party or lean towards the Democratic Party? [The Republican Party, The Democratic Party]

F.2.2 Tariff interview survey (June 2025)

(We elicit basic sociodemographic variables. We do not replicate these instructions here.)

In this study, you'll participate in an interview with a chatbot. This interview is a very important part of the study. Please answer thoughtfully and take your time. You have two options to answer the questions: You type the answers using your keyboard. You use the voice recording feature, see how it works in the video below. *[Short demonstration video]*

(New survey page)

Interview topic: On April 5, 2025, the United States began collecting a 10% tariff on all imports, with higher rates on some goods. On the next page, you will participate in an interview with an AI chatbot about your views about tariffs.

To ensure that the interview is as meaningful as possible, it's important that your responses provide sufficient detail. Please aim to write full sentences in your answers to allow for a meaningful conversation. There are no right or wrong answers. We expect the interview to take about 20 minutes to complete on average.

(New survey page)

(Chat interface for the interview)

(New survey page)

How would you rate your overall experience with the interview conducted by the AI chatbot? [Extremely positive, Very positive, Somewhat positive, Neutral, Somewhat negative, Very negative, Extremely negative]

How natural did the conversation with the AI chatbot feel? [Extremely natural, Very natural, Somewhat natural, Neutral, Somewhat unnatural, Very unnatural, Extremely unnatural]

If you were to participate in a future study involving a qualitative interview conducted through a chat interface, would you prefer texting with the same AI chatbot or an actual human interviewer? [Strongly prefer AI chatbot, Moderately prefer AI chatbot, Slightly prefer AI chatbot, Slightly prefer human interviewer, Moderately prefer human interviewer, Strongly prefer human interviewer]

Would you be interested in participating in an interview with an AI chatbot again? [Yes, No]

(New survey page)

How accurately did the previous interview capture how you personally think about the topic of tariffs? [Very accurate, Somewhat accurate, Neither accurate nor inaccurate, Somewhat inaccurate, Very inaccurate]

Did the interview questions seem neutrally phrased or did the questions already suggest a certain answer? [Very neutral, Somewhat neutral, Neither neutral nor suggestive, Somewhat suggestive, Very suggestive]

Were the interview questions open-ended, allowing you to answer in your own words, or closed-ended, limiting your response options? [Very open-ended, Somewhat open-ended, Neither open-ended nor closed-ended, Somewhat closed-ended, Very closed-ended]

How comfortable did you feel interacting with the AI chatbot? [Extremely comfortable, Very comfortable, Moderately comfortable, Slightly comfortable, Not at all comfortable]

(New survey page)

To what extent do you believe higher US tariffs will have positive or negative impacts on the overall US economy in the short and long term? [*Categories:* Short-term overall economic impact; Long-term overall economic impact]

[Response options for each element above: Very negative impact, Somewhat negative impact, Neither positive nor negative impact, Somewhat positive impact, Very positive impact]

(New survey page)

Over the next 12 months, how do you expect higher US tariffs to affect the following outcomes? [*Listed outcomes:* US inflation rate; US real GDP growth; US unemployment rate; US trade deficit; Net job creation in the US; US stock market returns (S&P 500); Variety of products available to US consumers; Global GDP growth]

[Response options for each element above: Strongly increase, Somewhat increase, Neither increase nor decrease, Somewhat decrease, Strongly decrease]

(New survey page)

To what extent do you expect each group to gain or lose from higher US tariffs? [*List of groups:* US manufacturers, US consumers, US exporters, High-income households, Low-income households]

[Response options for each group: Gain a lot, Gain a little, No effect, Lose a little, Lose a lot]

(New survey page)

How likely is it that major US trading partners will retaliate with their own tariffs?

[Very likely, Somewhat likely, Neither likely nor unlikely, Somewhat unlikely, Very unlikely]

Which statement best describes how firms will respond to higher tariffs on imported goods?

[Mostly absorb the higher cost themselves, Split the cost about 50/50 between firm and consumer, Mostly pass the cost on to US consumers]

(New survey page)

Please tell us whether you agree or disagree with the following statement: "In international trade, if one country makes more money, then it is generally the case that the other country makes less money."

[Strongly agree, Agree, Neither agree nor disagree, Disagree, Strongly disagree]

(New survey page)

In which of the following industries do you work in your main job? If you are currently unemployed, think about your last job. [Drop down: Constructions; Durable manufacturing: Finance, Insurance, and real estate; Health and education; Nondurable manufacturing; Retail and wholesale; Services; Transportation, communications and other public utilities; Other: ____]

How closely do you follow news about politics and the economy?

[Very closely, Somewhat closely, Not very closely, Not closely at all]

F.2.3 Tariff benchmark survey (July 2025)

(The survey is identical to the screener survey described in Section F.2.1 and the tariff interview study described in Section F.2.2, with the exception that we remove instructions specific to the interview. We replace the interview block with the instructions below.)

On April 5, 2025, the United States began collecting a 10% tariff on all imports, with higher rates on some goods. We will now ask you a few questions about your views on tariffs. Please aim to write full sentences. There are no right or wrong answers.

(New survey page. The following open-ended questions are shown on separate pages.)

What are the main reasons you support or oppose the new 10 percent tariff on imports?

Can you think of any potential benefits or costs of tariffs that you haven't mentioned so far?

Imagine you had to explain to someone unfamiliar with tariffs how they work and how they might affect different groups. What would you say?

Suppose the United States suddenly raised tariffs to 30 percent on imports from all major trading partners. What do you think would happen in the short run, and what might happen further down the road?

Is there anything else you think is important that we haven't touched on yet?

Is there anything you would want to know or see that could change your mind about the tariffs?